

Kubernetes

标签（空格分隔）： Kubernetes k8s

- Kubernetes
 - 标签（空格分隔）： Kubernetes k8s
 - 环境
 - Topology
 - SYS Init
 - Hostname And Hosts
 - Master
 - Docker CE
 - kubelet
 - Flannel
 - Node
 - Node Join Clusters
 - Admin Guide
 - K8s 系统架构
 - node
 - API Server
 - Scheduler
 - pod
 - Pod 生命周期
 - liveness
 - readiness
 - postStart preStop
 - control
 - Deployment
 - Daemonset
 - service
 - headless
 - sessionAffinity
 - Ingress Controller

- 提供 SSL 服务
 - Traefik
 - secret configmap
 - emptyDir
 - gitRepo
 - HostPath
 - NFS
 - volume pv pvc StorageClass
 - StatefulSet
 - rbac
 - role
 - serviceaccount
 - ResourceQuota
 - kubeconfig
 - CNI
 - ■ Flannel
 - Flannel Directrouting
 - Flannel host-gw
 - Calico
 - Dashboard
 - Monitor
 - Metrics-Server
 - customresourcedefinitions
 - Prometheus
 - HPA
 - helm
 - tiller
 - Gitlab Install By Helm
- Creating Highly Available Clusters with kubeadm
 - Init
 - Create load balancer for kube-apiserver
 - Etcd cluster
 - 先在 master01 生成 etcd 证书
 - 在所有的 master 都执行
 - 检查集群状态
 - 创建 k8s 集群

- Install CNI
 - Flannel And Enable Directrouting
 - Installing Calico for policy and flannel for networking
- Add Node
- Dashboard
 - 部署Web UI
 - 创建一个管理员用户
 - kubeconfig Login
- Error
- K8s With Paas
- Ingress
 - 启用 Ingress
 - gitlab 实例
 - 新加 node 并设置污点
 - gitlab安装
 - Traefik 设置
 - Ingress_nginx 设置
 - https

环境

```
CentOS Linux release 7.4.1708 (Core)
3.10.0-693.el7.x86_64
```

Topology

```
172.28.28.70 qa-k8s-master01.mljr.com
172.28.28.71 qa-k8s-master02.mljr.com
172.28.28.76 qa-k8s-master02.mljr.com
172.28.28.72 qa-k8s-node01.mljr.com
172.28.28.73 qa-k8s-node02.mljr.com
172.28.28.74 qa-k8s-node03.mljr.com
172.28.28.75 qa-k8s-node04.mljr.com
```

SYS Init

```
cat>> /etc/rc.loacl <<EOF
modprobe ip_vs
EOF

cat >> /etc/bashrc <<EOF
export HISTTIMEFORMAT="%Y-%m-%d-%H:%M:%S "
EOF

cat >> /etc/security/limits.conf <<EOF
* soft nofile 40960
* hard nofile 40960
EOF

cat > /usr/bin/safecron <<- 'EOF'
#!/bin/bash
# remove -r function from crontab
# PATH

if [[ "$*" =~ "-r" ]] ; then
    echo "SB! Dangerous..."
    echo "Exit..."
    exit 2
else
    /usr/bin/crontab $*
    exit 0
fi
EOF

chmod +x /usr/bin/safecron
grep "crontab='/usr/bin/safecron'" /etc/bashrc || echo "alias crontab='/usr/bin/safecron'" >> /etc/bashrc

setenforce 0
sed -i 's/SELINUX=enforcing/SELINUX=disabled/' /etc/selinux/config

systemctl stop firewalld.service
systemctl disable firewalld.service

systemctl stop iptables.service
systemctl disable iptables.service
```

```
swapoff -a
sed -i 's/.*swap.*/#&/' /etc/fstab

cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=http://mirrors.aliyun.com/kubernetes/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=0
repo_gpgcheck=0
gpgkey=http://mirrors.aliyun.com/kubernetes/yum/doc/yum-key.gpg
    http://mirrors.aliyun.com/kubernetes/yum/doc/rpm-package-key.gpg
EOF

cd /usr/local/src
wget http://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm
rpm -ivh epel-release-latest-7.noarch.rpm

crontab -e <<EOF
:1,1s#^#0 */2 * * * /usr/sbin/ntpdate 1.cn.pool.ntp.org \
:wq
EOF
```

Hostname And Hosts

```
# 每台服务服务器 各自执行

hostnamectl --pretty set-hostname qa-k8s-master03.mljr.com
hostnamectl --transient set-hostname qa-k8s-master03.mljr.com
hostnamectl --static set-hostname qa-k8s-master03.mljr.com
```

```
# hosts 文件添加
172.28.28.70 qa-k8s-master01.mljr.com master01
172.28.28.71 qa-k8s-master02.mljr.com master02
172.28.28.76 qa-k8s-master03.mljr.com master03
172.28.28.72 qa-k8s-node01.mljr.com node01
172.28.28.73 qa-k8s-node02.mljr.com node02
172.28.28.74 qa-k8s-node03.mljr.com node03
```

```
172.28.28.75 qa-k8s-node04.mljr.com node04
```

Master

Docker CE

```
yum remove docker docker-common docker-selinux docker-engine -y  
  
yum install -y yum-utils device-mapper-persistent-data lvm2  
  
yum-config-manager --add-repo https://download.docker.com/linux/centos/docker-ce.repo  
  
yum -y install docker-ce  
  
systemctl enable docker && systemctl start docker
```

kubelet

```
yum install -y kubelet kubeadm kubectl  
systemctl enable kubelet && systemctl start kubelet
```

禁止iptables对bridge数据进行处理

```
# 所有节点都要执行  
  
cat <<EOF > /etc/sysctl.d/k8s.conf  
net.bridge.bridge-nf-call-ip6tables = 1  
net.bridge.bridge-nf-call-iptables = 1  
EOF  
sysctl --system
```

master init

```
kubeadm init --kubernetes-version=v1.11.0 --pod-network-cidr=10.244.0.0/16 --apiserver-advertise-address=172.28.28.70 --ignore-preflight-errors=all
```

//正常执行结果如下

Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

You can now join any number of machines by running the following on each node as root:

```
kubeadm join 172.28.28.70:6443 --token en8yv2.njpxt99hcv6hctcf --discovery-token-ca-cert-hash sha256:3b800b8dc8c6bc23565fd80ffa143ef7e4659324f4c10706fad83d33b7bc0a
```

```
useradd k8s  
mkdir -p /home/k8s/.kube  
cp -i /etc/kubernetes/admin.conf /home/k8s/.kube/config  
chown k8s. /home/k8s/.kube/config  
su - k8s
```

在 k8s 用户下执行 kubectl 命令 如:

```
$ kubectl get nodes
```

Flannel

// 也可以用 Calico, 基于 BGP 协议, 比 Flannel 更强大

```
$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

Node

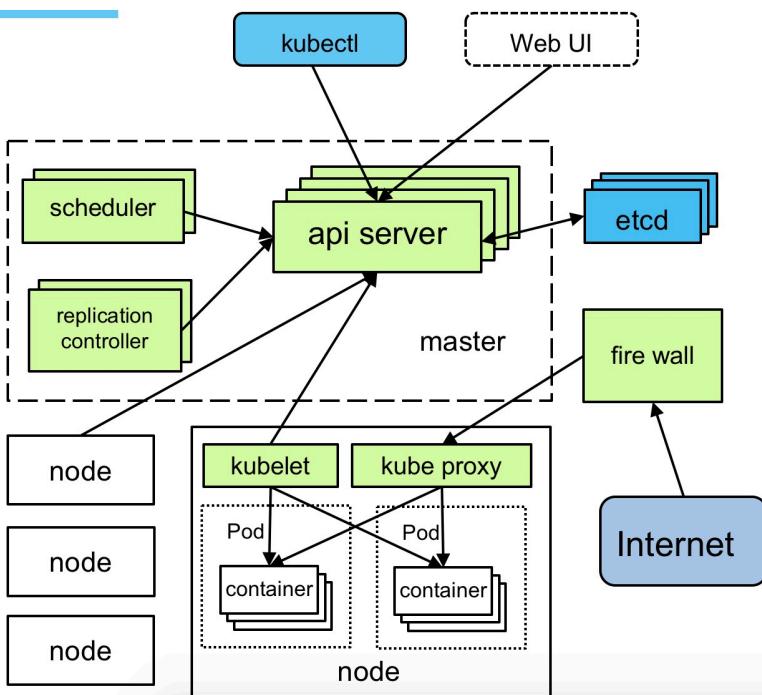
```
yum install docker-ce kubelet kubeadm kubectl -y  
  
systemctl enable docker && systemctl start docker  
systemctl enable kubelet.service
```

Node Join Clusters

```
kubeadm join 172.28.28.70:6443 --token en8yv2.njpxt99hcv6hctcf --discovery-token-ca-cert-hash sha256:3b800b8dc8c6bc23565fd80ffafe143ef7e4659324f4c10706fad83d3  
3b7bc0a
```

Admin Guide

K8s 系统架构



- 用户通过kubectl提交需要运行的docker container(pod)
- api server把请求存储在etcd里面
- scheduler扫描，分配机器
- kubelet找到自己需要跑的container，在本机上运行
- 用户提交RC描述，replication controller监视集群中的容器并保持数量
- 用户提交service描述文件，由kube proxy负责具体的工作流量转发

1) etcd

Etcd是一个分布式的高性能key-value存储系统，负责Kubernetes 中所有 REST API 对象的持久化保存。

2) kube-apiserver

Kubernetes 对外提供统一的入口，通过 REST API 的方式供客户端调用，例如 kubectl.

3) kube-controller-manager

控制集群中程序任务的后台线程。逻辑上，每个控制任务对应一个线程，包括**Node Controller**, Replication Controller, Service Controller, Endpoints Controller等。

4) kube-scheduler

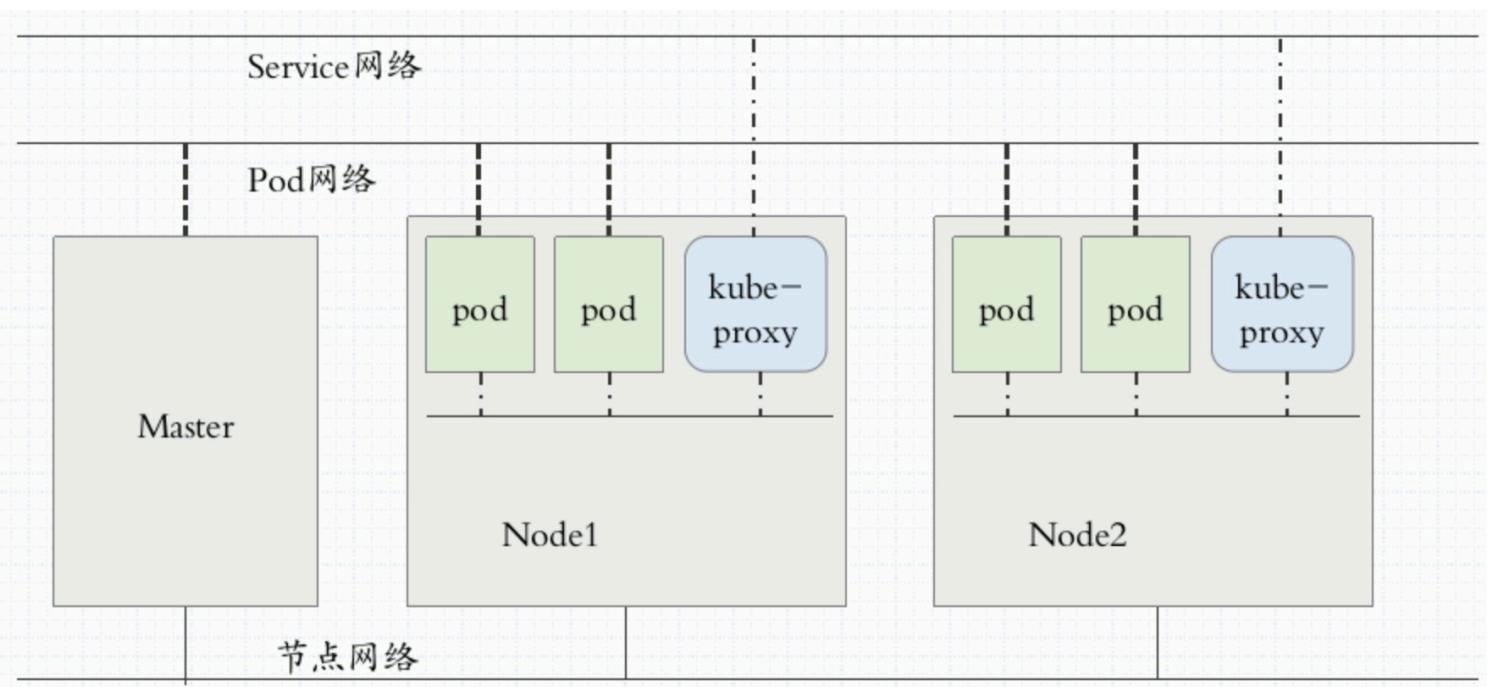
维护可用的 **node 节点及资源列表**，根据输入的待调度**pod**请求，通过调度算法选择最优**node节点**，输出**绑定了pod的node节点**.

5) kubelet

Kubelet 是 Kubernetes 集群中每个**node和Api-server**的连接点，负责容器和pod的实际管理工作。

6) kube-proxy

为pod提供代理服务，用来实现 kubernetes 的 service 机制。



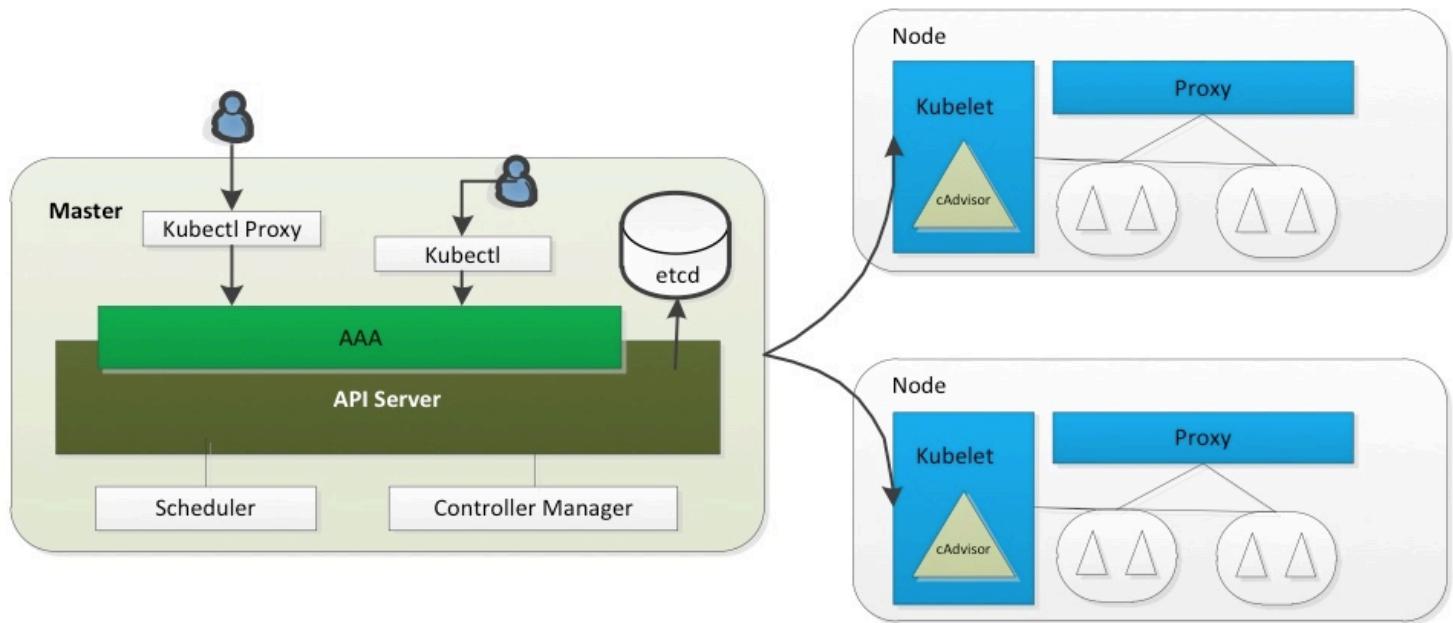
node

给 node 打标签，用于服务指定的服务。

```
kubectl label nodes qa-k8s-node01.mljr.com nodetype=vmware
```

```
[root@qa-k8s-master01 configmap]# cat pod-demo.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-nginx
  namespace: default
  labels:
    app: yyfq
    type: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.15-alpine
    imagePullPolicy: IfNotPresent
    ports:
    - name: http
      containerPort: 80
  nodeSelector:
    nodetype: vmware
```

API Server

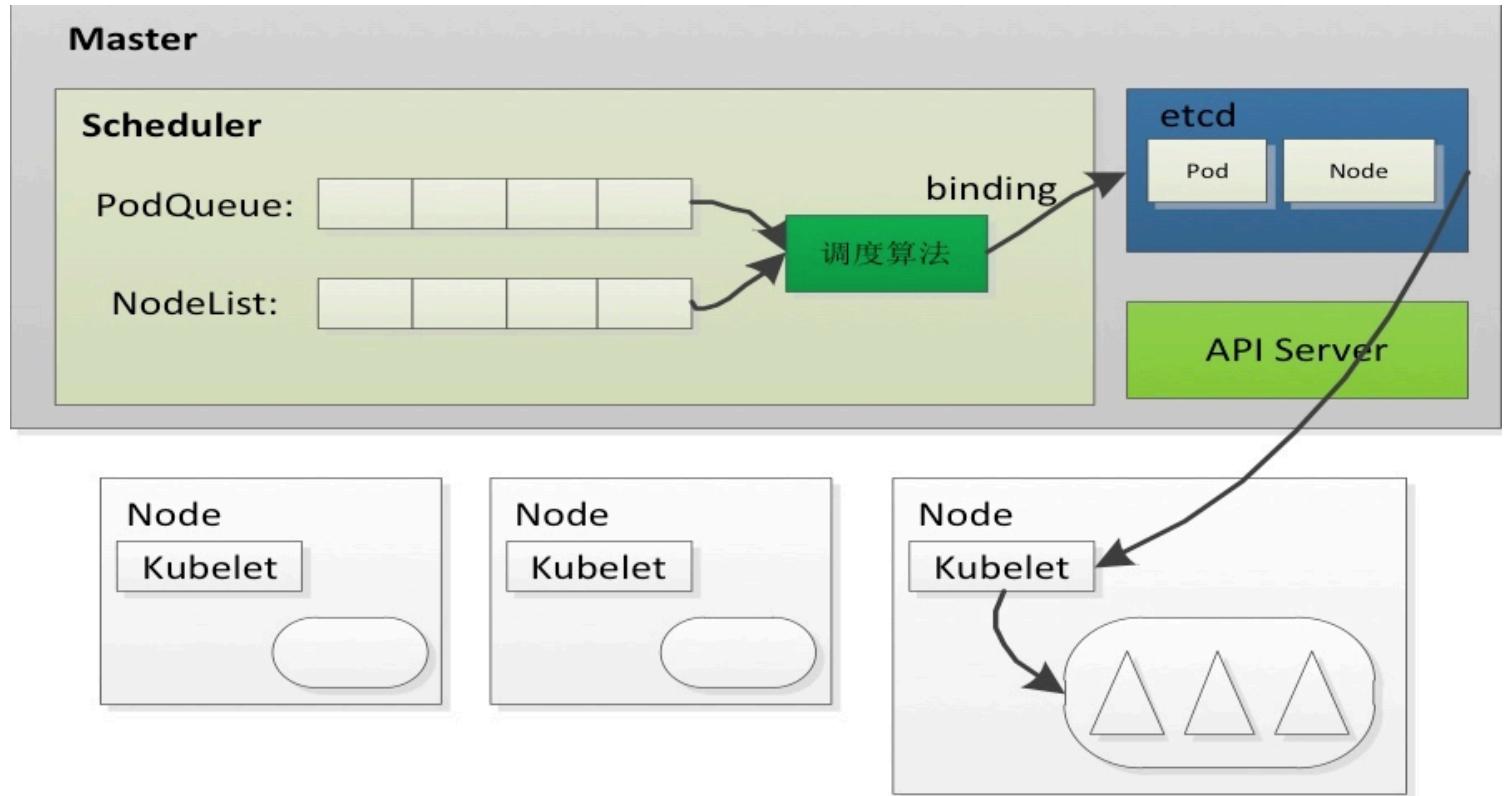


1) 提供集群管理的API接口；

- api-server进程提供[http\(8080\)](http://8080)/[https\(6443\)](https://6443)两个端口供访问;

2) 成为集群内各个功能模块之间数据交互和通信的中心枢纽

Scheduler



- 负责Pod调度
- 通过调度算法为待调度Pod列表的每个Pod从Node列表中选择一个最合适的Node

pod

Pod 是一组紧密关联的容器集合，它们共享 IPC, Network 和 UTS namespace，是 Kubernetes 调度的基本单位。

Pod 的设计理念是支持多个容器在一个 Pod 中共享网络和文件系统，可以通过进程间通信和文件共享这种简单高效的方式组合完成服务。

Pod 的特征：

包含多个共享 IPC, Network 和 UTC namespace 的容器，可直接通过 localhost 通信

所有 Pod 内容器都可以访问共享的 Volume，可以访问共享数据。

无容错性，直接创建的 Pod 一旦被调度后就跟 Node 绑定，即使 Node 挂掉也不会被重新调度（

而是被自动删除），因此推荐使用 Deployment, Daemonset 等控制器来容错

优雅终止，Pod 删除的时候先给其内的进程发送 SIGTERM，等待一段时间（grace period）后才强制停止依然还在运行的进程

特权容器（通过 SecurityContext 配置）具有改变系统配置的权限（在网络插件中大量应用）

```
// Kubernetes v1.8+ 还支持容器间共享 PID namespace，需要 docker >= 1.13.1，并配置 kubelet --docker-disable-shared-pid=false.
```

```
[root@qa-k8s-master01 configmap]# cat pod-demo.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-nginx
  namespace: default
  labels:
    app: yyfq
    type: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.15-alpine
    imagePullPolicy: IfNotPresent
    ports:
    - name: http
      containerPort: 80
```

```
[root@qa-k8s-master01 configmap]# kubectl apply -f pod-demo.yaml
pod/pod-nginx created
[root@qa-k8s-master01 configmap]# kubectl get pods -o wide --show-labels
NAME           READY   STATUS    RESTARTS   AGE     IP
pod-nginx      1/1     Running   0          23s    10.244.2.9
  qa-k8s-node02.mljr.com   app=yyfq,type=nginx
```

```
// -l 做标签过滤，显示有 app 标签（可以指定多个逗号分开）的 pod
```

```
[root@qa-k8s-master01 configmap]# kubectl get pods -o wide -l app
NAME           READY   STATUS    RESTARTS   AGE     IP
pod-nginx      1/1     Running   0          5m     10.244.2.9
  qa-k8s-node02.mljr.com
```

```
// -l 也可以做 等值/不等值 的标签选择器，多个标签为与关系
[root@qa-k8s-master01 configmap]# kubectl get pods -l type=nginx,release=stable
NAME        READY   STATUS    RESTARTS   AGE
pod-nginx   1/1     Running   0          16m
[root@qa-k8s-master01 configmap]# kubectl get pods -l type=nginx,release!=stable
No resources found.
[root@qa-k8s-master01 configmap]# kubectl get pods -l type=nginx,release
NAME        READY   STATUS    RESTARTS   AGE
pod-nginx   1/1     Running   0          16m

// -L 显示知道标签的值
[root@qa-k8s-master01 configmap]# kubectl get pods -L app,type
NAME          READY   STATUS    RESTARTS   AGE   APP
TYPE
pod-nginx     1/1     Running   0          7m   yyfq
nginx
[root@qa-k8s-master01 configmap]# kubectl get pods -l "release in (canary,beta,stable)"
NAME        READY   STATUS    RESTARTS   AGE
pod-nginx   1/1     Running   0          19m
[root@qa-k8s-master01 configmap]# kubectl get pods -l "release notin (canary,beta,stable)"
No resources found.
[root@qa-k8s-master01 configmap]# kubectl get pods -l "release in (canary,beta,alpha)"
No resources found.

// 给 pod 打标签
[root@qa-k8s-master01 configmap]# kubectl label --overwrite pods pod-nginx release=stable
pod/pod-nginx labeled
[root@qa-k8s-master01 configmap]# kubectl get pods -o wide --show-labels
NAME        READY   STATUS    RESTARTS   AGE   IP           LABELS
NODE
pod-nginx   1/1     Running   0          11m   10.244.2.9   app=yyfq,release=stable,type=nginx
```

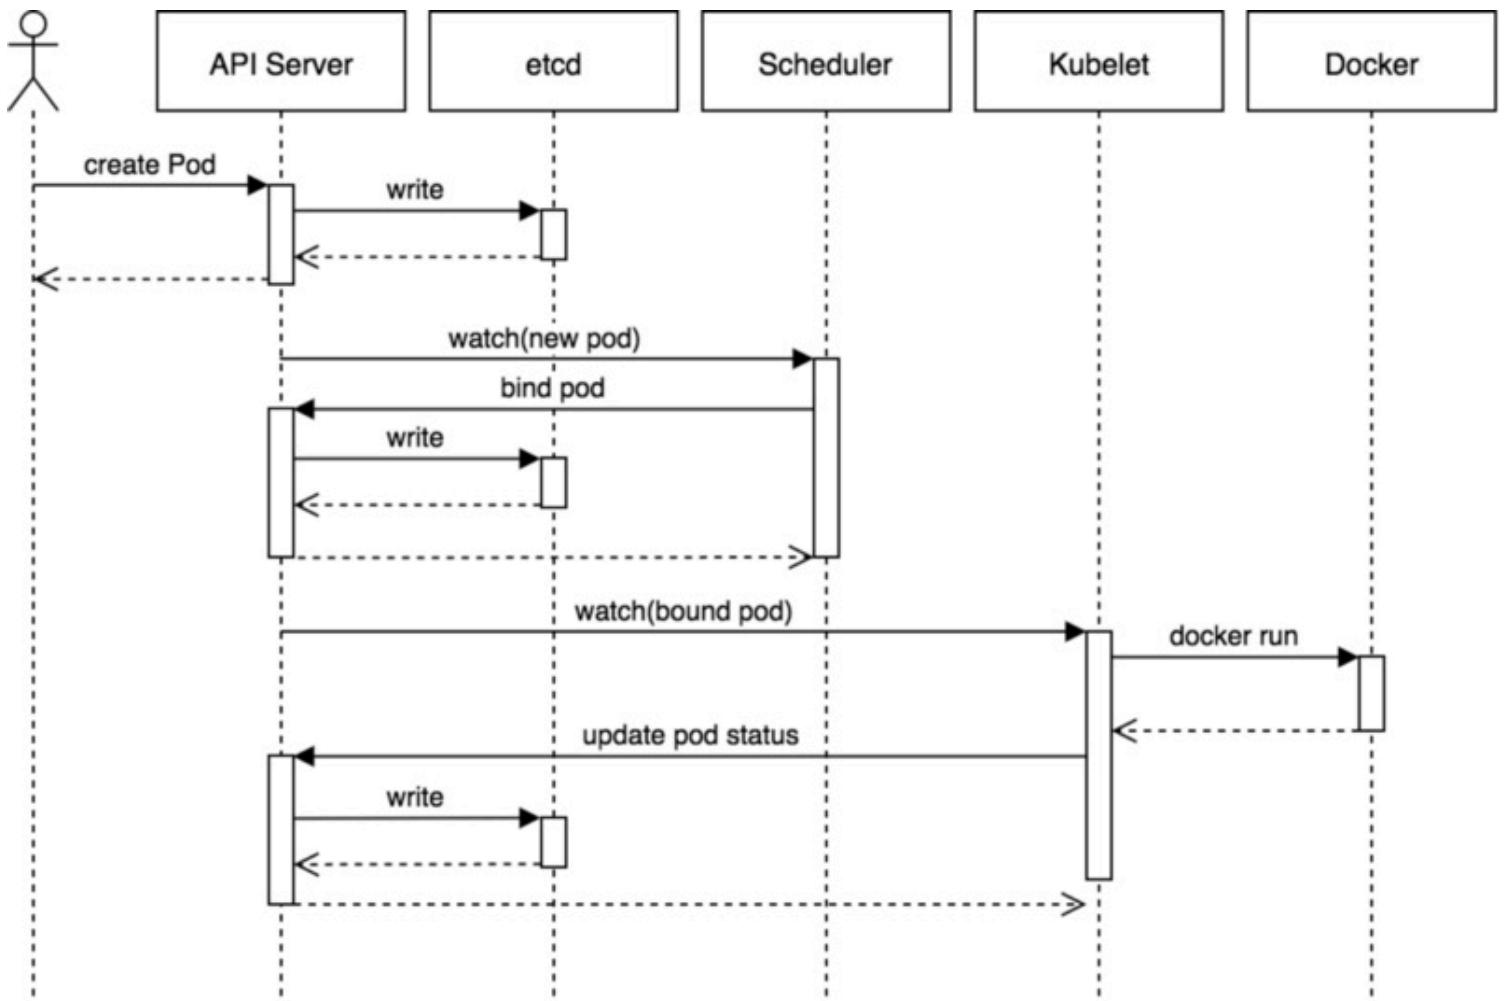
```
// 给 node 打标签
[root@qa-k8s-master01 configmap]# kubectl get node
NAME        STATUS   ROLES      AGE   VERSION
qa-k8s-master01.mljr.com  Ready    master    8d    v1.11.2
qa-k8s-node01.mljr.com   Ready    <none>   7d    v1.11.2
```

```

qa-k8s-node02.mljr.com     Ready    <none>    7d      v1.11.2
qa-k8s-node03.mljr.com     Ready    <none>    7d      v1.11.2
qa-k8s-node04.mljr.com     Ready    <none>    7d      v1.11.2
[root@qa-k8s-master01 configmap]# kubectl label nodes qa-k8s-node01.mljr.com nodetype=vmware
node/qa-k8s-node01.mljr.com labeled
[root@qa-k8s-master01 configmap]# kubectl get node --show-labels -l nodetype
NAME           STATUS   ROLES   AGE    VERSION   LABELS
qa-k8s-node01.mljr.com   Ready    <none>   7d    v1.11.2   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/fluentd-ds-ready=true,beta.kubernetes.io/os=linux,kubernetes.io/hostname=qa-k8s-node01.mljr.com,nodetype=vmware
[root@qa-k8s-master01 configmap]# cat pod-demo.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-nginx
  namespace: default
  labels:
    app: yyfq
    type: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.15-alpine
    imagePullPolicy: IfNotPresent
    ports:
    - name: http
      containerPort: 80
  nodeSelector:
    nodetype: vmware
[root@qa-k8s-master01 configmap]# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP
pod-nginx     1/1     Running   0          20s   10.244.1.13
qa-k8s-node01.mljr.com

```

Pod 生命周期



状态: Pending, Running, Failed, Succeeded, Unknown

Pod生命周期中的重要行为:

初始化容器

容器探测:

liveness: 存活性探测

readiness: 就绪性探测

pod 容器策略:

restartPolicy:

Always, OnFailure, Never. Default to Always.

liveness

```
[root@qa-k8s-master01 configmap]# cat liveness-exec-demo.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-liveness-exec
```

```

namespace: dev
labels:
  app: yyfq
  type: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.15-alpine
      imagePullPolicy: IfNotPresent
      ports:
        - name: http
          containerPort: 80
      command: ["/bin/sh","-c","touch /tmp/healthy;sleep 30;rm -f /tmp/healthy;sleep 1200"]

```

```

livenessProbe:
  exec:
    command: ["test","-c","/tmp/healthy"]
  initialDelaySeconds: 1
  periodSeconds: 3
  failureThreshold: 2
  timeoutSeconds: 1

```

```

restartPolicy: Always

```

```

[root@qa-k8s-master01 configmap]# kubectl apply -f liveness-exec-demo.yaml
pod/pod-liveness-exec created

```

```

[root@qa-k8s-master01 configmap]# kubectl get pods -n dev

```

NAME	READY	STATUS	RESTARTS	AGE
pod-liveness-exec	1/1	Running	0	17s

```

[root@qa-k8s-master01 configmap]# kubectl get pods -n dev -w

```

NAME	READY	STATUS	RESTARTS	AGE
pod-liveness-exec	0/1	CrashLoopBackOff	5	5m
pod-liveness-exec	1/1	Running	6	5m

// 可以观察得到，一直检查容器，是否正常，发现失败，一直重启。

```

// httpGet 探针
[root@qa-k8s-master01 configmap]# cat liveness-httpget-demo.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-liveness-httpget
  namespace: dev
  labels:
    app: yyfq

```

```

type: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.15-alpine
      imagePullPolicy: IfNotPresent
      ports:
        - name: http
          containerPort: 80
      livenessProbe:
        httpGet:
          path: /index.html
          port: http
        initialDelaySeconds: 1
        periodSeconds: 3
        failureThreshold: 2
        timeoutSeconds: 1
      restartPolicy: Always

// 注意此刻用的 dev 命名空间
// 可以连到容器，删除 index.html 进行测试

```

readiness

```

[root@qa-k8s-master01 configmap]# cat readiness-httpget-demo.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-readiness-httpget
  namespace: dev
  labels:
    app: yyfq
    type: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.15-alpine
      imagePullPolicy: IfNotPresent
      ports:
        - name: http
          containerPort: 80
      readinessProbe:
        httpGet:

```

```

    path: /index.html
    port: http
    initialDelaySeconds: 1
    periodSeconds: 3
    restartPolicy: Always
[root@qa-k8s-master01 configmap]# kubectl apply -f readiness-httpget-demo.yaml
pod/pod-readness-httpget created
[root@qa-k8s-master01 configmap]# kubectl get pods -n dev -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP           N
ODE
pod-liveness-httpget  1/1     Running   7          27m   10.244.2.10  q
a-k8s-node02.mljr.com
pod-readness-httpget  0/1     Running   0          3s    10.244.2.11  q
a-k8s-node02.mljr.com
[root@qa-k8s-master01 configmap]# kubectl get pods -n dev -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP           N
ODE
pod-liveness-httpget  1/1     Running   7          27m   10.244.2.10  q
a-k8s-node02.mljr.com
pod-readness-httpget  1/1     Running   0          7s    10.244.2.11  q
a-k8s-node02.mljr.com

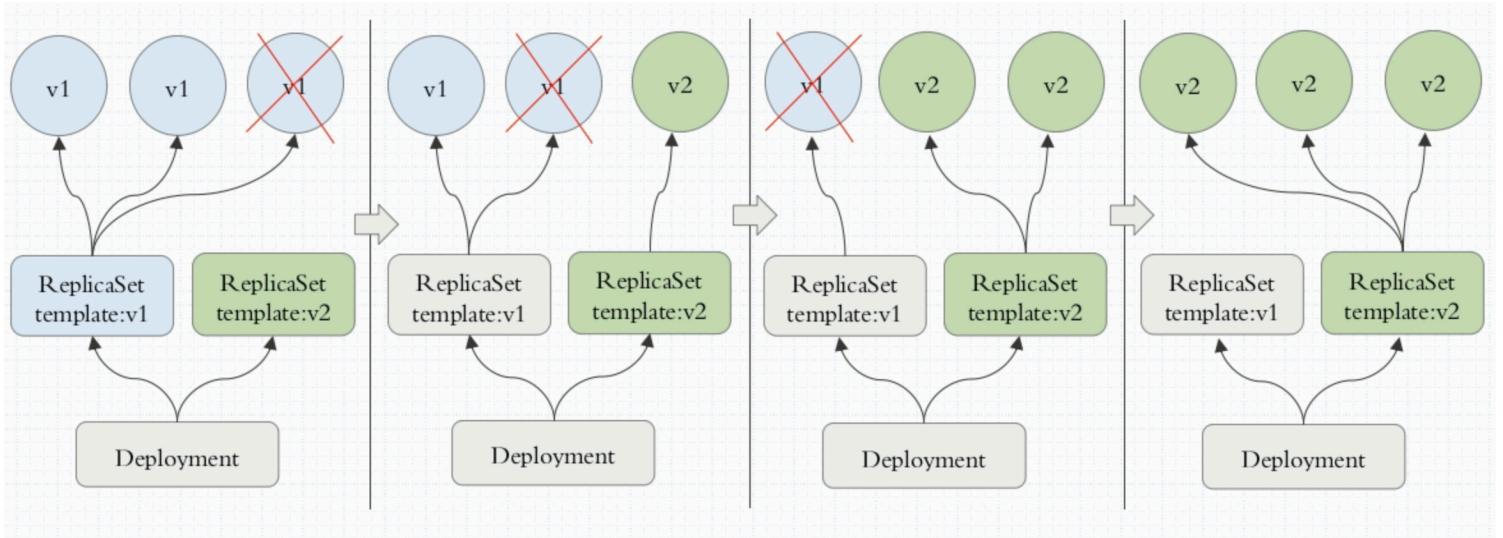
```

// 此时可以连上 pod-readness-httpget 容器，删除 index.html 文件，可以看出 READY 变成了 0/1.

postStart preStop

control

Deployment



```
[root@qa-k8s-master01 configmap]# cat deployment-demo.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-delploy
  namespace: dev
spec:
  replicas: 2
  selector:
    matchLabels:
      app: myapp
      release: canary
  template:
    metadata:
      labels:
        app: myapp
        release: canary
    spec:
      containers:
        - name: myapp
          image: nginx:1.15-alpine
          imagePullPolicy: IfNotPresent
        ports:
          - name: nginx
            containerPort: 80
```

```
[root@qa-k8s-master01 configmap]# kubectl apply -f deployment-demo.yaml
deployment.apps/myapp-delploy created
[root@qa-k8s-master01 configmap]# kubectl get deploy -n dev
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
myapp-delpoy	2	2	2	2	5m
[root@qa-k8s-master01 configmap]# kubectl get rs -n dev -o wide					
NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS
IMAGES					
myapp-delpoy-6b786474dc	2	2	2	2m	myapp
nginx:1.15-alpine app=myapp,pod-template-hash=2634203087,release=canary					
[root@qa-k8s-master01 configmap]# kubectl get pod -n dev					
NAME	READY	STATUS	RESTARTS	AGE	
myapp-delpoy-6b786474dc-jkj2w	1/1	Running	0	3m	
myapp-delpoy-6b786474dc-npnfr	1/1	Running	0	3m	

Daemonset

```
[root@qa-k8s-master01 configmap]# cat daemonset-demo.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis
  namespace: default
spec:
  replicas: 1
  selector:
    matchLabels:
      app: redis
      role: logstor
  template:
    metadata:
      labels:
        app: redis
        role: logstor
    spec:
      containers:
        - name: redis
          image: redis:4.0-alpine
          imagePullPolicy: IfNotPresent
        ports:
          - name: redis
            containerPort: 6379
```

```
---
```

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
```

```

name: filebeat-ds
namespace: default
spec:
  selector:
    matchLabels:
      app: filebeat
      release: stable
  template:
    metadata:
      labels:
        app: filebeat
        release: stable
    spec:
      containers:
        - name: filebeat
          image: ikubernetes/filebeat:5.6.5-alpine
          imagePullPolicy: IfNotPresent
          env:
            - name: REDIS_HOST
              value: redis.default.svc.cluster.local
            - name: REDIS_LOG_LEVEL
              value: info
[root@qa-k8s-master01 configmap]# kubectl apply -f daemonset-demo.yaml
deployment.apps/redis created
daemonset.apps/filebeat-ds created
// 暴露 redis 端口
// kubectl expose deployment redis --port=6379
[root@qa-k8s-master01 configmap]# kubectl expose deployment redis --port=6379
service/redis exposed

// 4个 node 节点
[root@qa-k8s-master01 configmap]# kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
filebeat-ds-92wcq   1/1     Running   1          15m
filebeat-ds-nzlg9    1/1     Running   0          15m
filebeat-ds-qrn5x    1/1     Running   0          15m
filebeat-ds-xbtvp    1/1     Running   0          15m
redis-5b5d6fbdd-4ljq   1/1     Running   0          16m

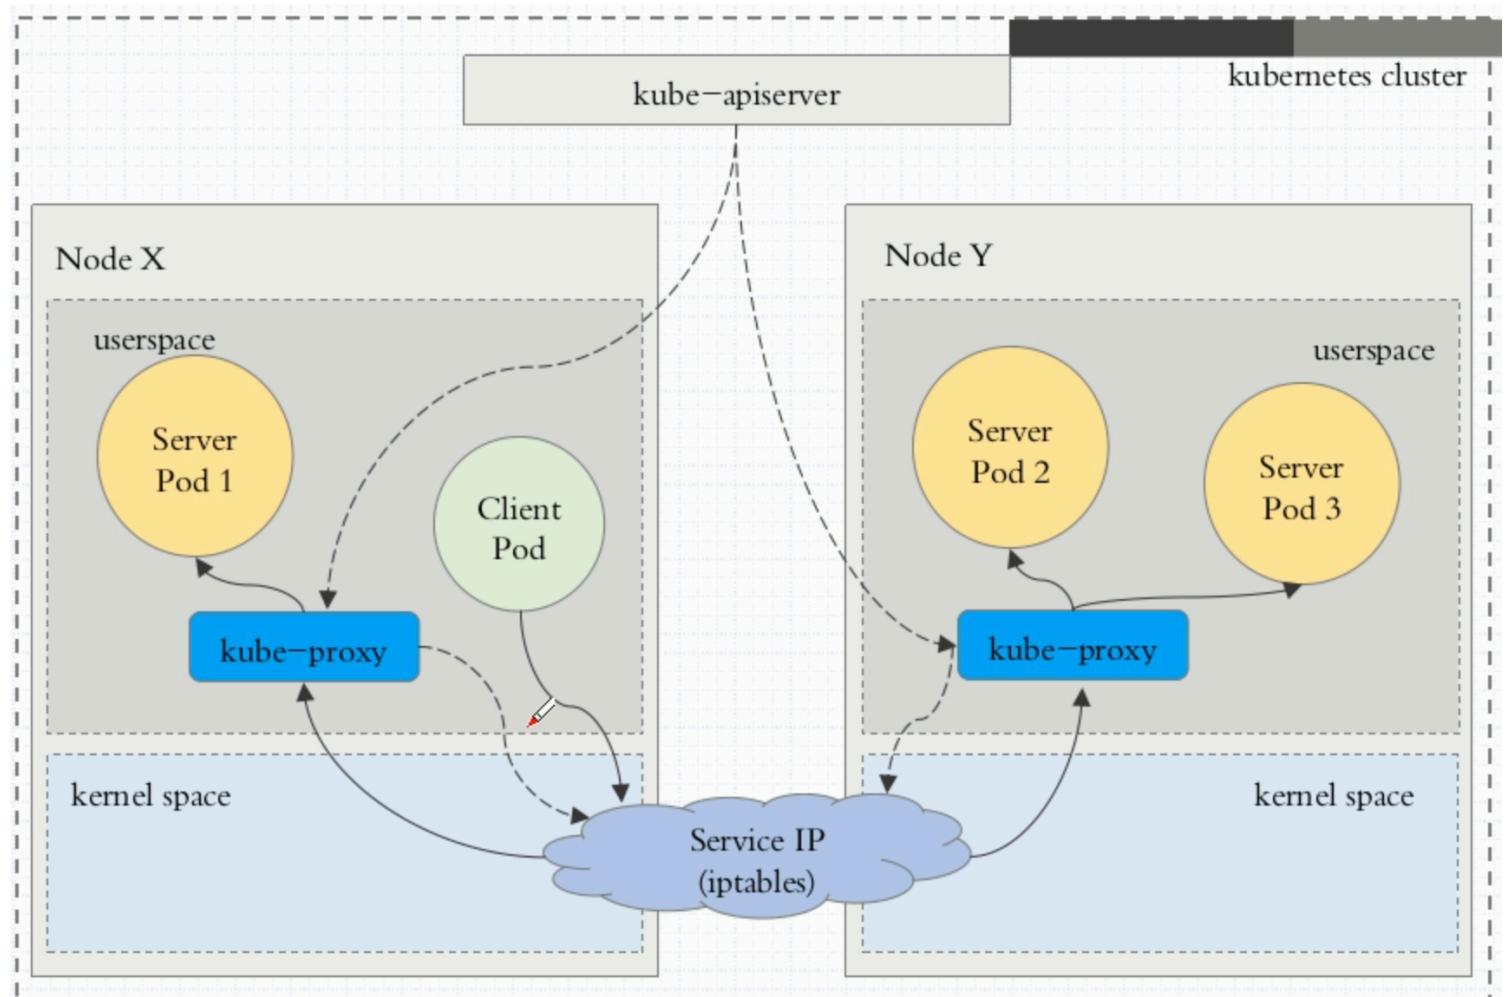
```

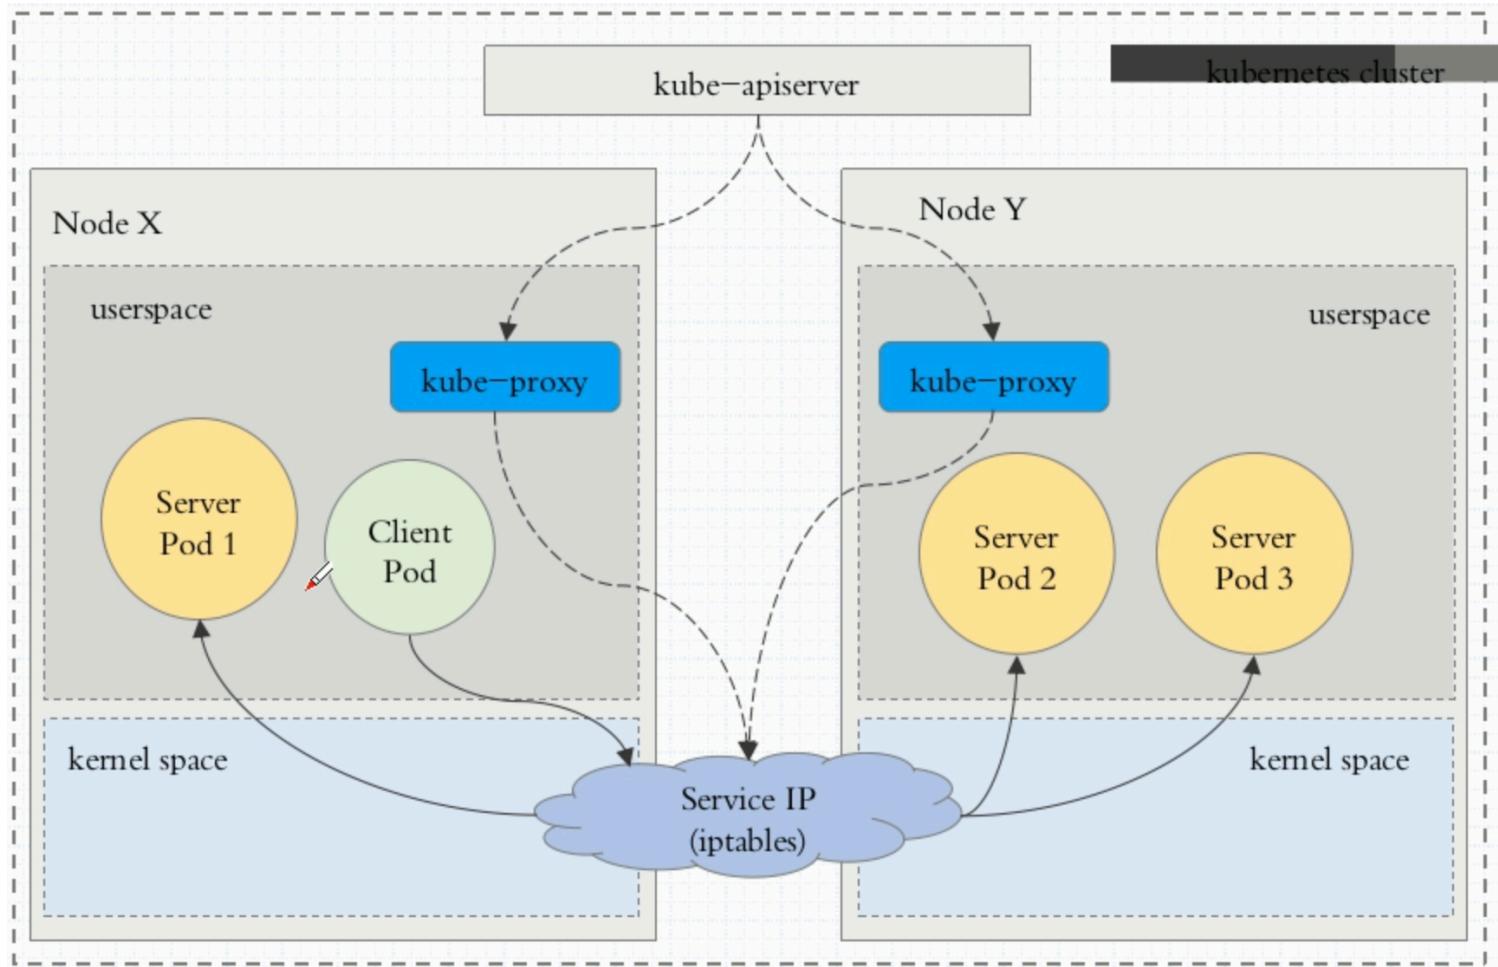
service

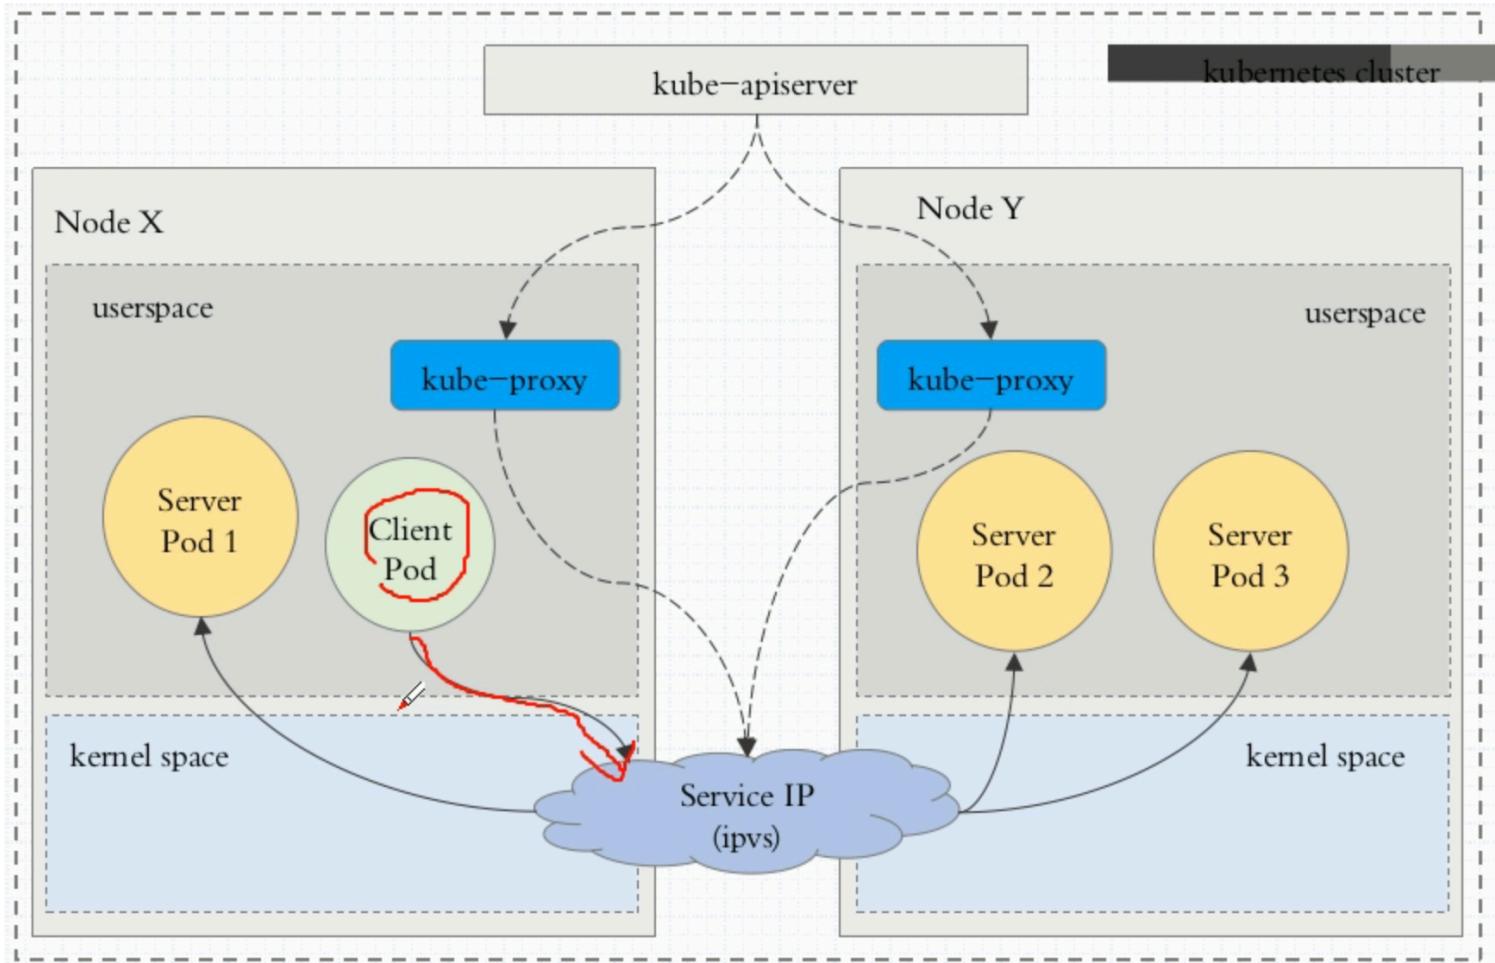
一个 **Kubernetes** 的 **Service** 是一种抽象，它定义了一个**Pod**的逻辑集合和一个用于访问它们的策略
- 有的时候称之为微服务。

一个Service的目标Pod集合通常是由Label Selector 来决定的.

kube-proxy 中实现 service 的方案: userspace, iptables 和 ipvs.







clusterip 例子

```
[root@qa-k8s-master01 configmap]# kubectl apply -f svc-clusterip-demo.yaml
service/redis created
deployment.apps/redis created
[root@qa-k8s-master01 configmap]# kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
pod-nginx             1/1     Running   0          1d
pod-nginx-hostpath   1/1     Running   0          2d
pod-nginx-pvc         1/1     Running   0          1d
pod-nginx-statefulset-0 1/1     Running   0          1d
pod-nginx-statefulset-1 1/1     Running   0          1d
pod-nginx-statefulset-2 1/1     Running   0          1d
redis-5b5d6fbabd-kt6dk 1/1     Running   0          26s
[root@qa-k8s-master01 configmap]# kubectl get svc
NAME        TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kubernetes  ClusterIP  10.96.0.1   <none>       443/TCP   9d
myapp       ClusterIP  None        <none>       80/TCP    1d
redis       ClusterIP  10.97.97.98  <none>       6379/TCP  35s
[root@qa-k8s-master01 configmap]# kubectl get deploy
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE

```

```
redis 1 1 1 43s
```

NodePort 例子

```
[root@qa-k8s-master01 configmap]# kubectl apply -f svc-nodeport-demo.yaml
service/nginx created
deployment.apps/myapp-delploy created
[root@qa-k8s-master01 configmap]# kubectl get pods -n dev
NAME READY STATUS RESTARTS AGE
myapp-delploy-b6bc574fb-kmvmz 1/1 Running 0 9s
myapp-delploy-b6bc574fb-rl8xs 1/1 Running 0 9s
[root@qa-k8s-master01 configmap]# kubectl -n dev get svc
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
nginx NodePort 10.100.91.108 <none> 80:31659/TCP 34s
[root@qa-k8s-master01 configmap]# kubectl -n dev get deploy
NAME DESIRED CURRENT UP-TO-DATE AVAILABLE AGE
myapp-delploy 2 2 2 2 39s
```

externalName

headless

有时你不需要一个单独的服务IP地址，也不需要做负载均衡。

在这种情况下，你可以创建一个“**headless**”的Service，只需要把集群IP(spec.clusterIP)指定为“**None**”即可。

对于这种类型的Service，没有集群IP地址的分配。也不会有DNS的配置来为一个Service名称返回多个A记录，这些记录会直接指向支撑这个Service的Pod。

另外，kube-proxy不会处理这些service，并且平台不会为它们做负载均衡或者代理。

endpoints controller仍然会在API中创建Endpoints的记录。

```
[root@qa-k8s-master01 configmap]# kubectl apply -f svc-headless-demo.yaml
service/nginx-svc created
deployment.apps/nginx-delploy created
[root@qa-k8s-master01 configmap]# kubectl get svc
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 9d
nginx-svc ClusterIP None <none> 80/TCP 18s
```

```
[root@qa-k8s-master01 configmap]# dig -t A nginx-svc.default.svc.cluster.local.
@10.96.0.10
```

```
; <>> DiG 9.9.4-RedHat-9.9.4-50.el7 <>> -t A nginx-svc.default.svc.cluster.local. @10.96.0.10
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 65232
;; flags: qr aa rd ra; QUERY: 1, ANSWER: 2, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;nginx-svc.default.svc.cluster.local. IN A

;; ANSWER SECTION:
nginx-svc.default.svc.cluster.local. 5 IN A 10.244.1.29
nginx-svc.default.svc.cluster.local. 5 IN A 10.244.2.22

;; Query time: 0 msec
;; SERVER: 10.96.0.10#53(10.96.0.10)
;; WHEN: 四 8月 23 00:49:01 CST 2018
;; MSG SIZE rcvd: 166
```

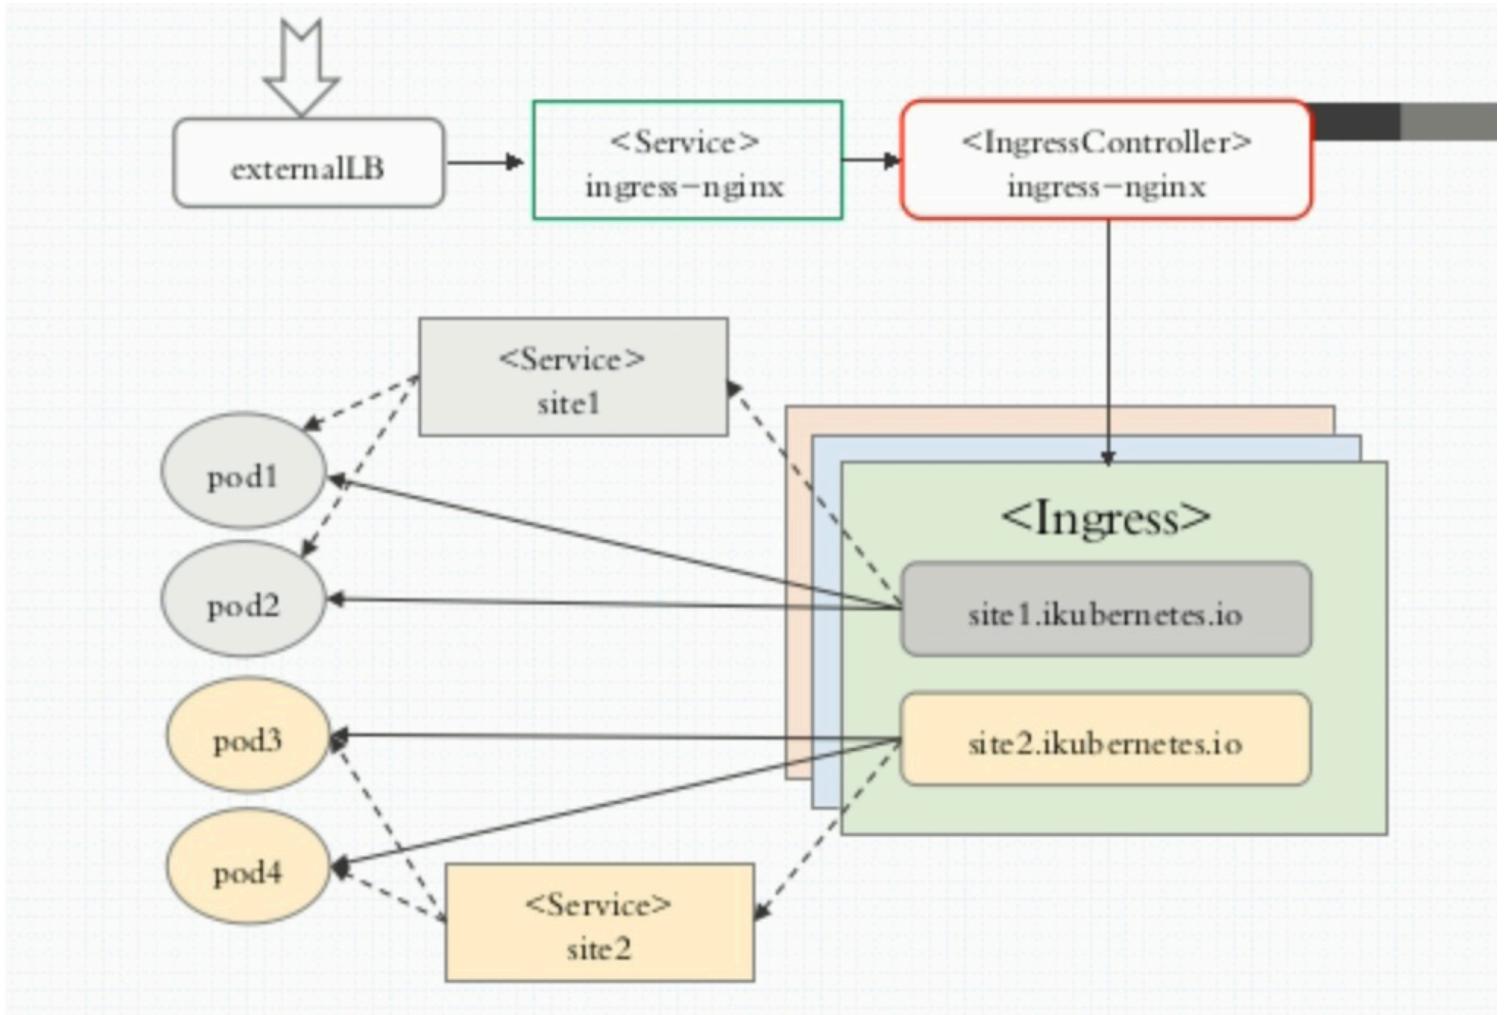
sessionAffinity

```
[root@qa-k8s-master01 configmap]# kubectl patch svc nginx-svc -p '{"spec":{"sessionAffinity":"ClientIP"}}'
service/nginx-svc patched
```

Ingress Controller

Ingress Controller 是一个或者一组提供7层代理的 pod.

Ingress 资源，可以基于后端 **server**(仅仅做后端资源收集，不做调用) 监控收集到后端 pod，后端 pod 资源改变后，Ingress可以直接编辑 Ingress Controller，



```
// 安裝 Ingress Nginx Controller https://kubernetes.github.io/ingress-nginx/
kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/master/deploy/mandatory.yaml
```

准备一个 NodePort Service

```
[root@qa-k8s-master01 ingress]# cat service-nodeport.yaml
apiVersion: v1
kind: Service
metadata:
  name: ingress-nginx
  namespace: ingress-nginx
  labels:
    app.kubernetes.io/name: ingress-nginx
    app.kubernetes.io/part-of: ingress-nginx
spec:
  type: NodePort
```

```

ports:
- name: http
  port: 80
  targetPort: 80
  nodePort: 30080
  protocol: TCP
- name: https
  port: 443
  targetPort: 443
  nodePort: 30443
  protocol: TCP
selector:
  app.kubernetes.io/name: ingress-nginx

```

```

[root@qa-k8s-master01 ingress]# kubectl apply -f service-nodeport.yaml
service/ingress-nginx created
[root@qa-k8s-master01 ingress]# kubectl get svc -n ingress-nginx
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)
              AGE
default-http-backend   ClusterIP  10.99.176.38 <none>        80/TCP
                      4m
ingress-nginx       NodePort   10.106.101.154 <none>        80:30080/TCP,
443:30443/TCP      5s

```

准备 Ingress Server

```

[root@qa-k8s-master01 ingress]# cat ingress-tomcat-demo.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-tomcat
  namespace: default
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  rules:
  - host: tomcat.mljr.com
    http:
      paths:
      - path:
          backend:
            serviceName: tomcat
            servicePort: 8080

```

```
[root@qa-k8s-master01 ingress]# kubectl apply -f ingress-tomcat-demo.yaml
ingress.extensions/ingress-tomcat created
[root@qa-k8s-master01 ingress]# kubectl get ing
NAME          HOSTS           ADDRESS      PORTS      AGE
ingress-tomcat tomcat.mljr.com        80          25s
```

准备后端 tomcat service

```
[root@qa-k8s-master01 ingress]# cat deploy-tomcat-demo.yaml
apiVersion: v1
kind: Service
metadata:
  name: tomcat
  namespace: default
spec:
  selector:
    app: tomcat
    release: canary
  ports:
    - name: http
      targetPort: 8080
      port: 8080
    - name: ajp
      targetPort: 8009
      port: 8009
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: tomcat-deploy
  namespace: default
spec:
  replicas: 3
  selector:
    matchLabels:
      app: tomcat
      release: canary
  template:
    metadata:
      labels:
        app: tomcat
        release: canary
  spec:
    containers:
```

```
- name: tomcat
  image: tomcat:8.5.32-jre8-alpine
  ports:
    - name: http
      containerPort: 8080
    - name: ajp
      containerPort: 8009
```

```
[root@qa-k8s-master01 ingress]# kubectl apply -f deploy-tomcat-demo.yaml
service/tomcat created
```

```
deployment.apps/tomcat-deploy created
```

```
[root@qa-k8s-master01 ingress]# kubectl get pods -o wide
```

NAME	NODE	READY	STATUS	RESTARTS	AGE	IP
tomcat-deploy-588c79d48d-5qqbq	4.4.11 qa-k8s-node04.mljr.com	1/1	Running	0	58s	10.24
tomcat-deploy-588c79d48d-d8nhc	4.1.31 qa-k8s-node01.mljr.com	1/1	Running	0	58s	10.24
tomcat-deploy-588c79d48d-q4dst	4.2.23 qa-k8s-node02.mljr.com	1/1	Running	0	58s	10.24

提供 SSL 服务

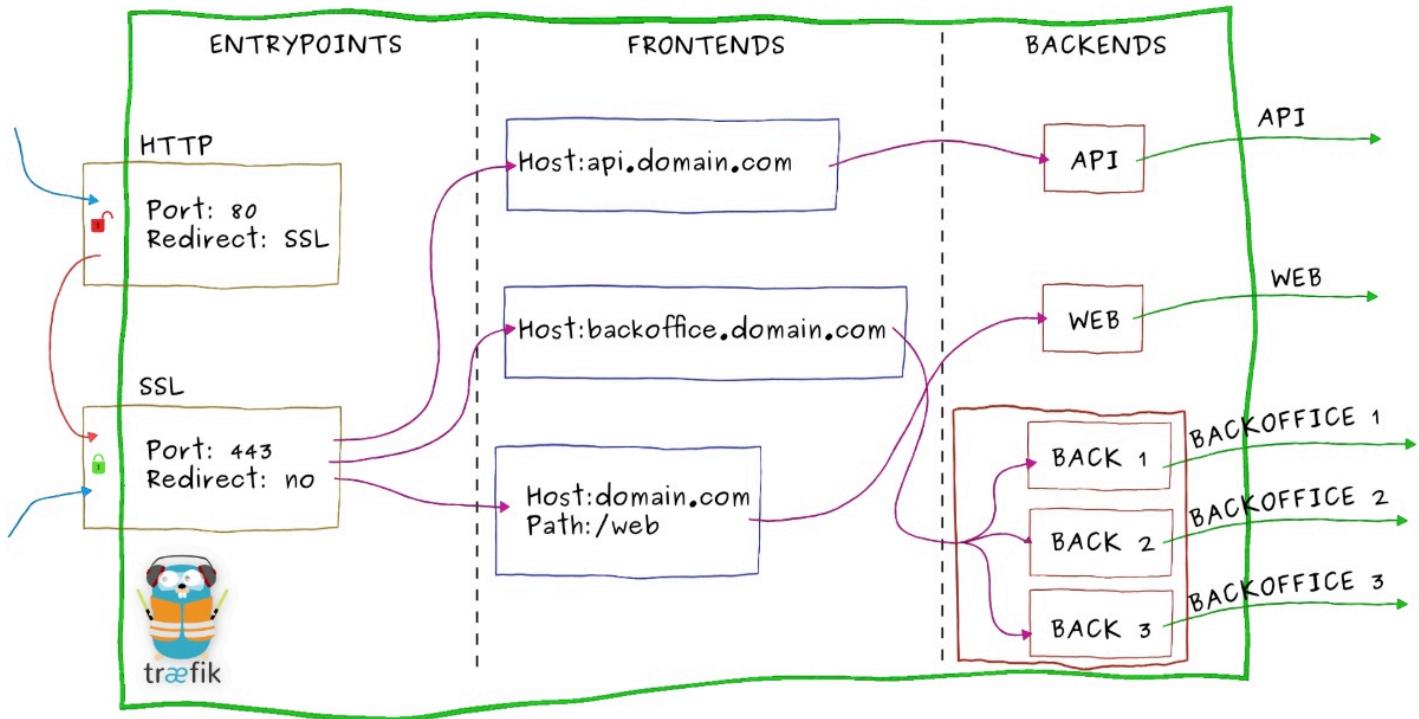
待补充

<http://docs.traefik.cn/>

Traefik

Ingress Controller 的存在就是为了能跟 kubernetes 交互，又能写 nginx 配置，还能 reload 它，这是一种折中方案。

而最近开始出现的 traefik 天生就是提供了对 kubernetes 的支持，也就是说 traefik 本身就能跟 kubernetes API 交互，感知后端变化，因此可以得知，在使用 traefik 时，整体架构如下。



secret configmap

配置容器化应用的方式：

1 自定义命令行参数

```
args: []
```

2 把配置文件直接熔进镜像

3 环境变量

(1) Cloud Native 的应用程序一般可直接通过环境变量加载配置

(2) 通过entrypoint脚本来预处理变量为配置文件中的配置信息

4 存储卷

```
// 配置文件从镜像解耦
configmap
secret
```

```
kubectl create configmap nginx-config --from-literal=nginx_port=8080
```

```
[root@qa-k8s-master01 ~]# kubectl get cm
NAME          DATA   AGE
nginx-config  1      25s
[root@qa-k8s-master01 ~]# kubectl describe cm nginx-config
Name:            nginx-config
Namespace:       default
Labels:          <none>
Annotations:    <none>

Data
=====
nginx_port:
-----
8080
Events: <none>
```

```
// 从文件
kubectl create cm nginx-www --from-file=./www.config

[root@qa-k8s-master01 configmap]# cat www.config
server {
    server_name myapp.mljr.com;
    listen 80;
    root /data/web/html;
}
[root@qa-k8s-master01 configmap]# kubectl get cm nginx-www -o yaml
apiVersion: v1
data:
  www.config: |
    server {
        server_name myapp.mljr.com;
        listen 80;
        root /data/web/html;
    }
kind: ConfigMap
metadata:
  creationTimestamp: 2018-08-20T07:46:43Z
  name: nginx-www
  namespace: default
  resourceVersion: "981539"
  selfLink: /api/v1/namespaces/default/configmaps/nginx-www
  uid: 2e555bd8-a44d-11e8-8cfc-0050569d5ecc
[root@qa-k8s-master01 configmap]#
```

```
[root@qa-k8s-master01 configmap]# cat pod1-configmap-demo.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-nginx-configmap
  namespace: default
spec:
  containers:
  - name: mynginx
    image: nginx:1.15-alpine
    imagePullPolicy: IfNotPresent
    ports:
    - name: mynginx
      containerPort: 80
    env:
    - name: NGINX_SERVER_PORT
      valueFrom:
        configMapKeyRef:
          name: nginx-config
          key: nginx_port
    volumeMounts:
    - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
      name: default-token-6chr7
      readOnly: true
  tolerations:
  - effect: NoExecute
    key: node.kubernetes.io/not-ready
    operator: Exists
    tolerationSeconds: 300
  - effect: NoExecute
    key: node.kubernetes.io/unreachable
    operator: Exists
    tolerationSeconds: 300
  volumes:
  - name: default-token-6chr7
    secret:
      defaultMode: 420
      secretName: default-token-6chr7
```

```
[root@qa-k8s-master01 configmap]# kubectl apply -f pod1-configmap-demo.yaml
pod/pod-nginx-configmap created
[root@qa-k8s-master01 configmap]# kubectl exec -it pod-nginx-configmap -- /bin/
sh
/ #
```

```
/ # env
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_SERVICE_PORT=443
HOSTNAME=pod-nginx-configmap
SHLVL=1
HOME=/root
NGINX_SERVER_PORT=8080
TERM=xterm
NGINX_VERSION=1.15.2
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_SERVICE_PORT_HTTPS=443
PWD=/
KUBERNETES_SERVICE_HOST=10.96.0.1
/ # exit
[root@qa-k8s-master01 configmap]#
```

emptyDir

```
// 定义 volume
// 在容器上 挂载 存储器 volumMounts
// 创建 pv
// 定义好 存储类
//
```

gitRepo

gitRepo 建立在 emptyDir 基础

git 仓库发生改变， pod 不会自己改变。

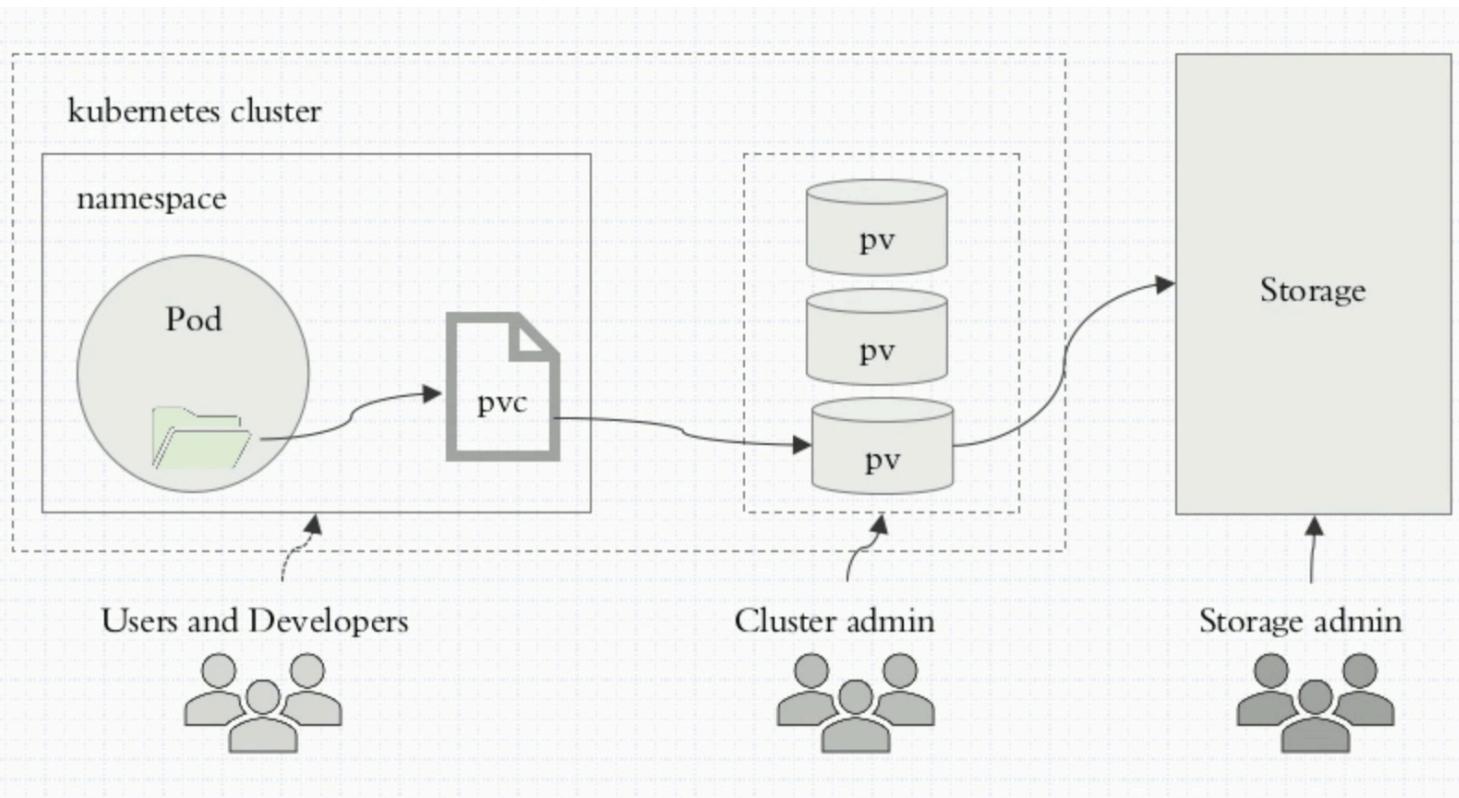
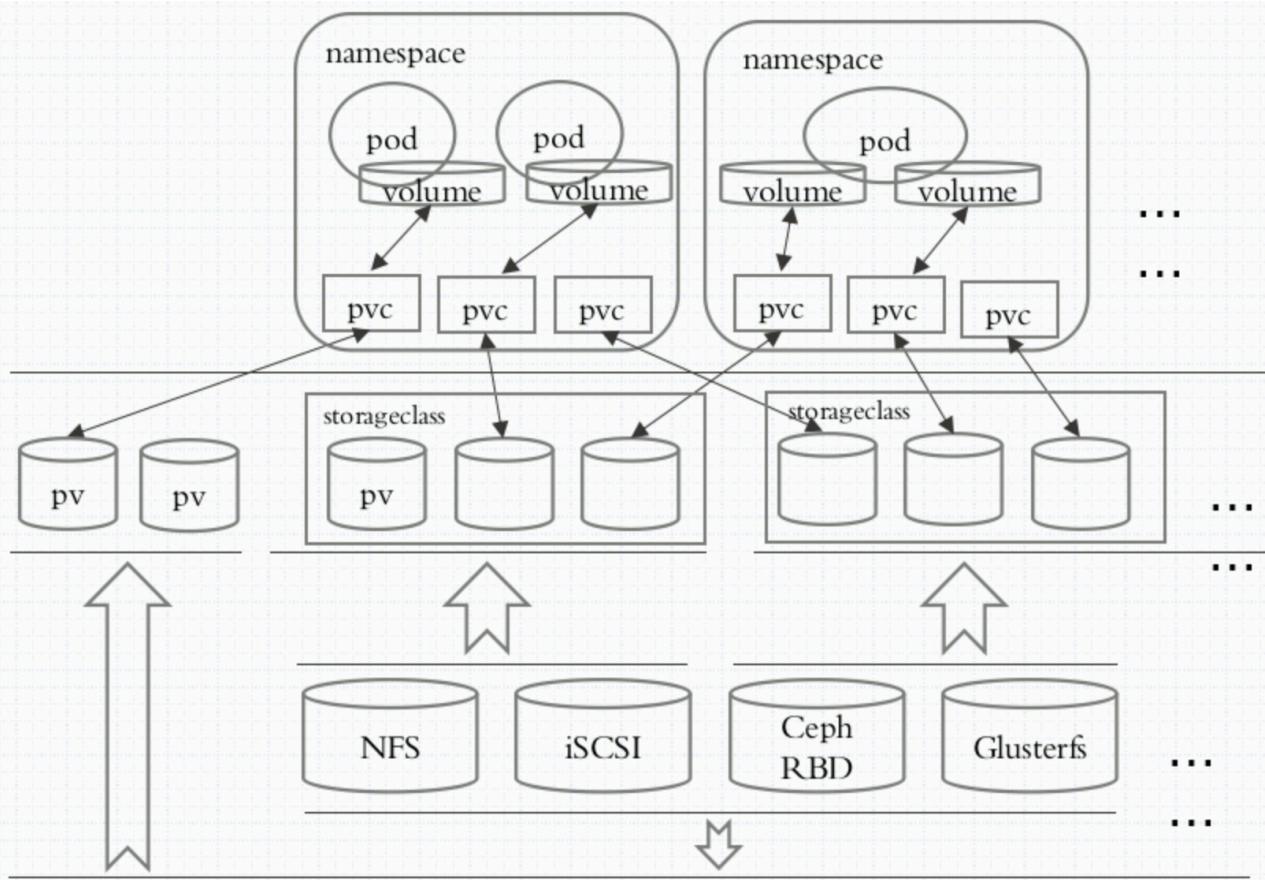
HostPath

```
// 节点级别的 持久存储  
  
// 所有节点  
mkdir -p /data/pod/volume1  
echo "$(hostname)" /data/pod/volume1/index.html
```

NFS

```
// 所有节点安装  
yum install nfs-utils -y  
  
// nfs server  
mkdir -p /pv/volumes  
[root@qa-k8s-master03 ~]# cat /etc/exports  
/pv/volumes 172.0.0.0/8(rw,no_root_squash)  
[root@qa-k8s-master03 ~]# systemctl start nfs  
  
// master 节点  
[root@qa-k8s-master01 configmap]# cat pod-nfs-demo.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: pod-nginx-nfs  
  namespace: default  
spec:  
  containers:  
    - name: nginx  
      image: nginx:1.15-alpine  
      imagePullPolicy: IfNotPresent  
      volumeMounts:  
        - name: html  
          mountPath: /usr/share/nginx/  
  volumes:  
    - name: html  
      nfs:  
        path: /pv/volumes  
        server: 172.28.28.76
```

volume pv pvc StorageClass



// k8s 工程师 创建 pvc, pv

```
mkdir -p /pv/volumes/v{1,2,3,4,5}
```

```
[root@qa-k8s-master03 volumes]# cat /etc/exports
/pv/volumes/v1 172.0.0.0/8(rw,no_root_squash)
/pv/volumes/v2 172.0.0.0/8(rw,no_root_squash)
/pv/volumes/v3 172.0.0.0/8(rw,no_root_squash)
/pv/volumes/v4 172.0.0.0/8(rw,no_root_squash)
/pv/volumes/v5 172.0.0.0/8(rw,no_root_squash)
[root@qa-k8s-master03 volumes]# systemctl daemon-reload
[root@qa-k8s-master03 volumes]# systemctl restart nfs
[root@qa-k8s-master03 volumes]# exportfs -arv
exporting 172.0.0.0/8:/pv/volumes/v5
exporting 172.0.0.0/8:/pv/volumes/v4
exporting 172.0.0.0/8:/pv/volumes/v3
exporting 172.0.0.0/8:/pv/volumes/v2
exporting 172.0.0.0/8:/pv/volumes/v1
[root@qa-k8s-master03 volumes]# showmount -e
Export list for qa-k8s-master03.mljr.com:
/pv/volumes/v5 172.0.0.0/8
/pv/volumes/v4 172.0.0.0/8
/pv/volumes/v3 172.0.0.0/8
/pv/volumes/v2 172.0.0.0/8
/pv/volumes/v1 172.0.0.0/8
```

```
[root@qa-k8s-master01 configmap]# cat pod-vol-pvc.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mypvc
  namespace: default
spec:
  accessModes: ["ReadWriteMany", "ReadWriteOnce"]
  resources:
    requests:
      storage: 6Gi
---
apiVersion: v1
kind: Pod
metadata:
  name: pod-nginx-pvc
  namespace: default
spec:
  containers:
  - name: nginx
    image: nginx:1.15-alpine
```

```

imagePullPolicy: IfNotPresent
volumeMounts:
- name: html
  mountPath: /usr/share/nginx/
volumes:
- name: html
  persistentVolumeClaim:
    claimName: mypvc
[root@qa-k8s-master01 configmap]# kubectl get pv
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS     CLAIM
STORAGECLASS   REASON     AGE
pv001      5Gi        RWO,RWX       Retain          Available
                                         10m
pv002      5Gi        RWO,RWX       Retain          Available
                                         10m
pv003      10Gi       RWO,RWX      Retain          Bound      default/mypvc
                                         10m
pv004      10Gi       RWO,RWX      Retain          Available
                                         10m
pv005      20Gi       RWO,RWX      Retain          Available
                                         12m

```

StatefulSet

volumeClaimTemplates 能自动完成两个功能：

- 为 pod 定义 volume
- 在 pod 名称空间中 自动创建 pvc

```

[root@qa-k8s-master01 configmap]# kubectl get pv
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS     CLAIM
STORAGECLASS   REASON     AGE
pv001      5Gi        RWO,RWX       Retain          Bound      default/html-po
d-nginx-statefulset-0
                                         15h
pv002      5Gi        RWO,RWX       Retain          Bound      default/html-po
d-nginx-statefulset-1
                                         15h
pv003      10Gi       RWO,RWX      Retain          Released   default/mypvc
                                         15h
pv004      10Gi       RWO,RWX      Retain          Bound      default/mypvc
                                         15h
pv005      20Gi       RWO,RWX      Retain          Bound      default/varlog
                                         15h

```

```
[root@qa-k8s-master01 configmap]# kubectl get pvc
```

NAME	ORAGECLASS	AGE	STATUS	VOLUME	CAPACITY	ACCESS MODES	ST
html-pod-nginx-statefulset-0		10m	Bound	pv001	5Gi	RwO,RwX	
html-pod-nginx-statefulset-1		10m	Bound	pv002	5Gi	RwO,RwX	
mypvc		15h	Bound	pv004	10Gi	RwO,RwX	
varlog		38m	Bound	pv005	20Gi	RwO,RwX	

```
[root@qa-k8s-master01 configmap]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
pod-nginx-hostpath	1/1	Running	0	16h
pod-nginx-pvc	1/1	Running	0	15h
pod-nginx-statefulset-0	1/1	Running	0	10m
pod-nginx-statefulset-1	1/1	Running	0	10m

```
[root@qa-k8s-master01 configmap]# kubectl delete -f pod-statefulset-demo.yaml
service "myapp" deleted
statefulset.apps "pod-nginx-statefulset" deleted
```

```
[root@qa-k8s-master01 configmap]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
pod-nginx-hostpath	1/1	Running	0	16h
pod-nginx-pvc	1/1	Running	0	15h
pod-nginx-statefulset-1	0/1	Terminating	0	12m

```
[root@qa-k8s-master01 configmap]# kubectl get pvc
```

NAME	ORAGECLASS	AGE	STATUS	VOLUME	CAPACITY	ACCESS MODES	ST
html-pod-nginx-statefulset-0		12m	Bound	pv001	5Gi	RwO,RwX	
html-pod-nginx-statefulset-1		12m	Bound	pv002	5Gi	RwO,RwX	
mypvc		15h	Bound	pv004	10Gi	RwO,RwX	
varlog		41m	Bound	pv005	20Gi	RwO,RwX	

```
[root@qa-k8s-master01 configmap]# kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
			STORAGECLASS		
pv001	5Gi	RwO,RwX	Retain	Bound	default/html-po
d-nginx-statefulset-0				15h	
pv002	5Gi	RwO,RwX	Retain	Bound	default/html-po
d-nginx-statefulset-1				15h	
pv003	10Gi	RwO,RwX	Retain	Released	default/mypvc
				15h	

pv004	10Gi	RW0, RWX	Retain	Bound 15h	default/mypvc
pv005	20Gi	RW0, RWX	Retain	Bound 15h	default/varlog
[root@qa-k8s-master01 configmap]# kubectl apply -f pod-statefulset-demo.yaml service/myapp created statefulset.apps/pod-nginx-statefulset created					

```
nslookup pod-nginx-statefulset-1.myapp.default.svc.cluster.local
```

// 字段如下

```
pod_name.service_name.namespce_name.svc.cluster.local
```

```
[root@qa-k8s-master01 configmap]# dig -t A pod-nginx-statefulset-1.myapp.default.svc.cluster.local @10.96.0.10
```

```
; <>> DiG 9.9.4-RedHat-9.9.4-50.el7 <>> -t A pod-nginx-statefulset-1.myapp.default.svc.cluster.local @10.96.0.10
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 1809
;; flags: qr aa rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;pod-nginx-statefulset-1.myapp.default.svc.cluster.local. IN A

;; ANSWER SECTION:
pod-nginx-statefulset-1.myapp.default.svc.cluster.local. 5 IN A 10.244.2.5

;; Query time: 0 msec
;; SERVER: 10.96.0.10#53(10.96.0.10)
;; WHEN: 二 8月 21 11:26:53 CST 2018
;; MSG SIZE rcvd: 155
```

// 扩容

```
kubectl scale statefulset pod-nginx-statefulset --replicas=3
```

```
[root@qa-k8s-master01 configmap]# kubectl scale statefulset pod-nginx-statefulset --replicas=3
statefulset.apps/pod-nginx-statefulset scaled
```

```
[root@qa-k8s-master01 configmap]# kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
pod-nginx-hostpath   1/1     Running   0          17h
pod-nginx-pvc        1/1     Running   0          15h
pod-nginx-statefulset-0 1/1     Running   0          23m
pod-nginx-statefulset-1 1/1     Running   0          23m
pod-nginx-statefulset-2 1/1     Running   0          29s
```

// 更新策略

```
[root@qa-k8s-master01 configmap]# kubectl explain sts.spec.updateStrategy.rollingUpdate
```

KIND: StatefulSet

VERSION: apps/v1

RESOURCE: rollingUpdate <Object>

DESCRIPTION:

RollingUpdate **is** used **to** communicate **parameters** when **Type is** RollingUpdateStatefulSetStrategyType.

RollingUpdateStatefulSetStrategy **is** used **to** communicate parameter **for** RollingUpdateStatefulSetStrategyType.

FIELDS:

partition <integer>

Partition indicates the ordinal **at** which the StatefulSet should be partitioned. **Default value is 0.**

partition : N 用于模拟金丝雀更新

表示 大于等于 N 的，才会更新。0 表示都更新。

// 查看 sts 更新策略

```
[root@qa-k8s-master01 configmap]# kubectl describe sts pod-nginx-statefulset | grep Update
```

Update Strategy: RollingUpdate

// 实现模拟金丝雀更新

```
[root@qa-k8s-master01 configmap]# kubectl patch sts pod-nginx-statefulset -p '{"spec":{"updateStrategy":{"rollingUpdate":{"partition":2}}}}'
```

// 然后更新测试

```
[root@qa-k8s-master01 configmap]# kubectl set image sts/pod-nginx-statefulset myapp=nginx:1.14-alpine
```

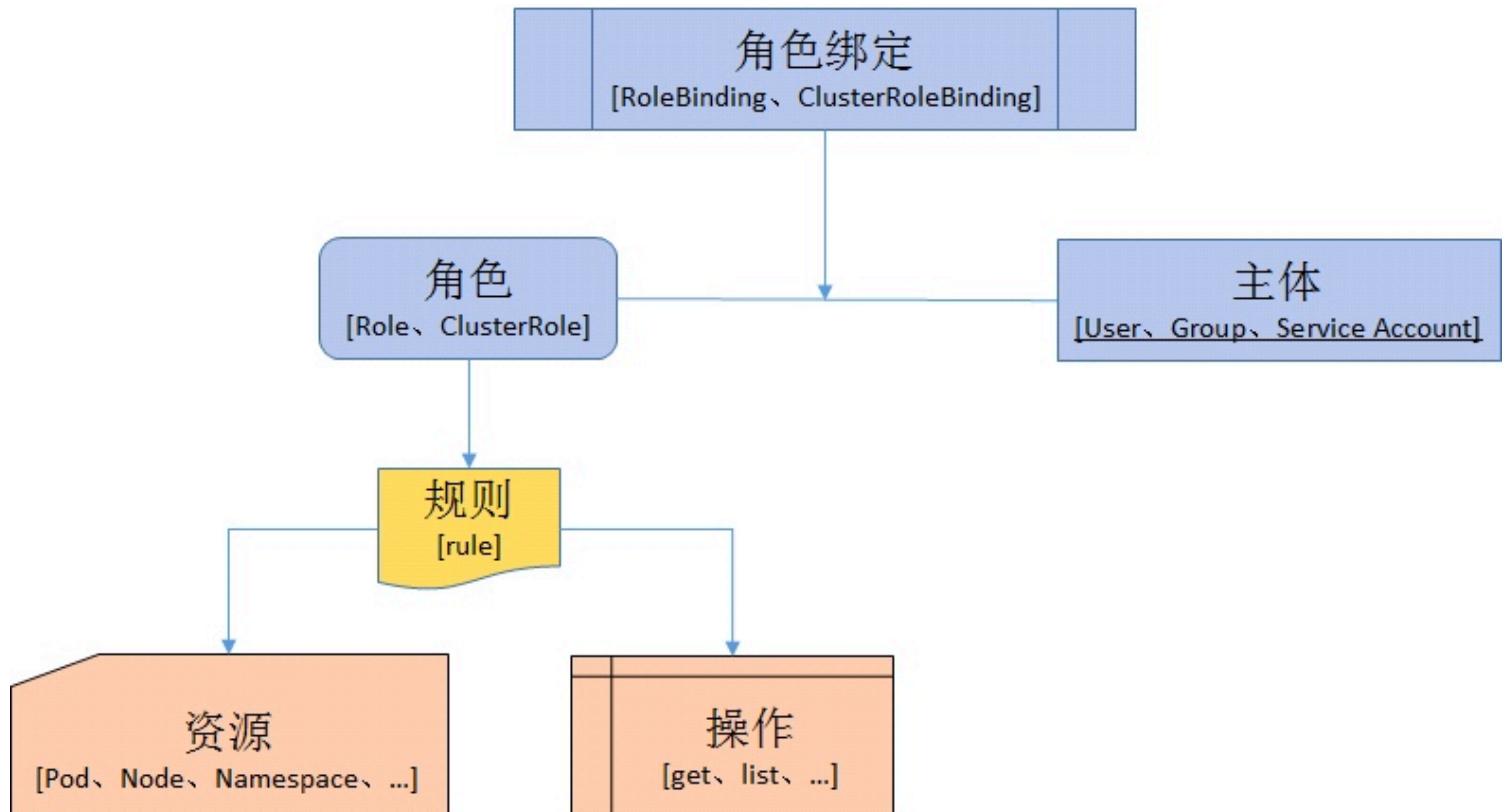
statefulset.apps/pod-nginx-statefulset image updated

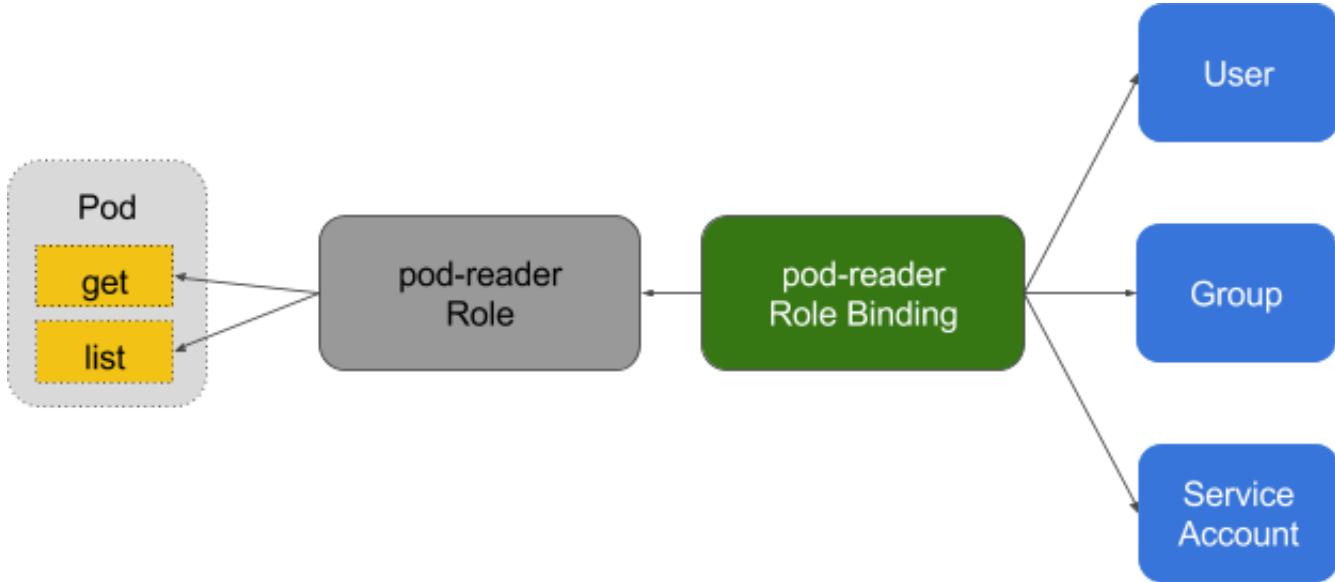
```
[root@qa-k8s-master01 configmap]# kubectl get sts -o wide
```

NAME	DESIRED	CURRENT	AGE	CONTAINERS	IMAGES
pod-nginx-statefulset	3	3	1h	myapp	nginx:1.14-alpine

```
[root@qa-k8s-master01 configmap]# kubectl describe pod pod-nginx-statefulset-2
| grep 'Image:'
  Image:          nginx:1.14-alpine
[root@qa-k8s-master01 configmap]# kubectl describe pod pod-nginx-statefulset-1
| grep 'Image:'
  Image:          nginx:1.15-alpine
[root@qa-k8s-master01 configmap]# kubectl patch sts pod-nginx-statefulset -p '{"spec":{"updateStrategy":{"rollingUpdate":{"partition":0}}}}'
statefulset.apps/pod-nginx-statefulset patched
[root@qa-k8s-master01 configmap]# kubectl describe pod pod-nginx-statefulset-1
| grep 'Image:'
  Image:          nginx:1.14-alpine
[root@qa-k8s-master01 configmap]# kubectl describe pod pod-nginx-statefulset-0
| grep 'Image:'
  Image:          nginx:1.14-alpine
```

rbac





role

```
kubectl create role pods-reader --verb=get,list,watch --resource=pods --dry-run -o yaml > role-demo.yaml
```

```
kubectl apply -f role-demo.yaml
```

```
[root@qa-k8s-master01 ~]# kubectl create rolebinding dashboard-read-pods-rolebinding --role=pods-reader --user=dashboard --dry-run -o yaml > rolebinding-demo.yaml
```

```
[root@qa-k8s-master01 ~]# cat rolebinding-demo.yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  creationTimestamp: null
  name: dashboard-read-pods-rolebinding
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: pods-reader
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: User
  name: dashboard
```

serviceaccount

```
[root@qa-k8s-master01 networkpolicy]# kubectl create serviceaccount mysa -o yaml --dry-run
apiVersion: v1
kind: ServiceAccount
metadata:
  creationTimestamp: null
  name: mysa
```

```
[root@qa-k8s-master01 networkpolicy]# kubectl create serviceaccount myadmin
serviceaccount/myadmin created
[root@qa-k8s-master01 networkpolicy]# kubectl get sa
NAME        SECRETS   AGE
def-ns-default  1        9h
default       1        6d
myadmin       1        12s
[root@qa-k8s-master01 networkpolicy]# kubectl get secret
NAME          TYPE           DATA   AGE
E
def-ns-default-token-x4w2d  kubernetes.io/service-account-token  3      9h
default-token-6chr7         kubernetes.io/service-account-token  3      6d
myadmin-token-qpmvj        kubernetes.io/service-account-token  3      45s
[root@qa-k8s-master01 networkpolicy]#
```

```
kubectl get pods client-7c9999bd74-cqz6h -o yaml
```

ResourceQuota

kubeconfig

```
[root@qa-k8s-master01 kubernetes]# kubectl config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
    server: https://172.28.28.70:6443
```

```
name: kubernetes
contexts:
- context:
  cluster: kubernetes
  user: kubernetes-admin
  name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED
```

```
kubectl config set-credentials dashboard-admin --client-certificate=./dashboard.crt --client-key=./dashboard.key
```

```
kubectl config set-context dashboard-admin@kubernetes --cluster=kubernetes --user=dashboard-admin
```

```
kubectl config use-context dashboard-admin@kubernetes
```

CNI

Kubernetes 网络通信：

- (1) 容器间通信：同一个**Pod**内的多个容器间的通信，**lo**
- (2) **Pod**通信：**Pod IP <--> Pod IP**
- (3) **Pod与Service**通信：**PodIP <--> ClusterIP**
- (4) **Service**与集群外部客户端的通信；

Flannel

Flannel Directrouting

使用 Flannel，添加一台新机器到集群时，Flannel 做了三件事：

- 使用 etcd 为新增机器分配一个子网
- 在宿主机上创建一个虚拟桥接接口(名为 docker0 的桥接器)
- 设置一个网络转发后端

flannel 支持多种后端：

VxLAN: 创建一个虚拟的 VXLAN 接口

(1) vxlan

(2) Directrouting (同一二层网络使用 host-gw, 跨三层使用 vxlan)

host-gw: 通过远程机器 IP, 创建到子网的IP路由. 这要求运行 flannel 的不同主机在二层直接互通.

UDP:

```
[root@qa-k8s-master01 configmap]# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE     IP
NODE                               NOMINATED NODE
nginx-delpoy-869b48bf9-glqf8      1/1     Running   0          2m     10.244.4.94   qa-k8s-node03.mljr.com
nginx-delpoy-869b48bf9-ldvzm      1/1     Running   0          2m     10.244.3.5   qa-k8s-node04.mljr.com
```

```
[root@qa-k8s-node04 ~]# ip route show
default via 172.28.3.254 dev ens160 proto static metric 100
10.244.0.0/24 via 10.244.0.0 dev flannel.1 onlink
10.244.1.0/24 via 10.244.1.0 dev flannel.1 onlink
10.244.2.0/24 via 10.244.2.0 dev flannel.1 onlink
10.244.3.0/24 dev cni0 proto kernel scope link src 10.244.3.1
10.244.4.0/24 via 10.244.4.0 dev flannel.1 onlink
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1
172.28.3.0/24 dev ens160 proto kernel scope link src 172.28.3.55 metric 100
```

```
[root@qa-k8s-node03 ~]# ip route show
default via 172.28.3.254 dev ens160 proto static metric 100
10.244.0.0/24 via 10.244.0.0 dev flannel.1 onlink
10.244.1.0/24 via 10.244.1.0 dev flannel.1 onlink
10.244.2.0/24 via 10.244.2.0 dev flannel.1 onlink
10.244.3.0/24 via 10.244.3.0 dev flannel.1 onlink
10.244.4.0/24 dev cni0 proto kernel scope link src 10.244.4.1
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1
172.28.3.254 dev ens160 proto static scope link metric 100
172.28.28.0/24 dev ens160 proto kernel scope link src 172.28.28.74 metric 100
```

// 下载 kube-flannel.yml , 修改如下字段, 然后重新应用.

```
net-conf.json: |
{
  "Network": "10.244.0.0/16",
  "Backend": {
```

```

        "Type": "vxlan",
        "Directrouting": true
    }
}

```

```
kubectl delete -f kube-flannel.yml
kubectl apply -f kube-flannel.yml
```

// 重新生成pod节点查看路由

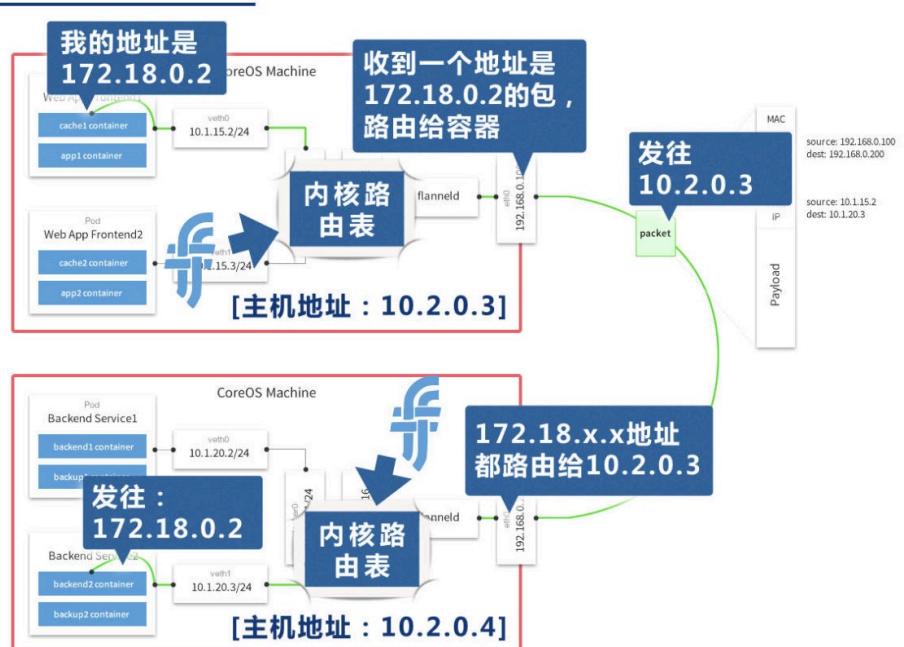
```
[root@qa-k8s-node03 ~]# ip route show
default via 172.28.3.254 dev ens160 proto static metric 100
10.244.0.0/24 via 172.28.28.75 dev ens160
10.244.1.0/24 via 172.28.28.71 dev ens160
10.244.2.0/24 via 172.28.28.76 dev ens160
10.244.3.0/24 via 10.244.3.0 dev flannel.1 onlink
10.244.4.0/24 dev cni0 proto kernel scope link src 10.244.4.1
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1
172.28.3.254 dev ens160 proto static scope link metric 100
172.28.28.0/24 dev ens160 proto kernel scope link src 172.28.28.74 metric 100
```

Flannel host-gw

问题二：容器地址不可达

Flannel的解决方案：
方案2. 主机路由

对应host-gw运行模式



局限性：

- 必须二层网络可达
- 适用于小规模集群，或是专用的大二层网络环境

Calico

Installing Calico for policy and flannel for networking

```
kubectl apply -f \
https://docs.projectcalico.org/v3.1/getting-started/kubernetes/installation/hosted/canal/rbac.yaml
```

```
kubectl apply -f \
https://docs.projectcalico.org/v3.1/getting-started/kubernetes/installation/hosted/canal/canal.yaml
```

```
mkdir /etc/kubernetes/networkpolicy
cd /etc/kubernetes/networkpolicy
```

```
// 建立 2个 命名空间 dev prod
kubectl create ns dev
kubectl create ns prod
```

```
// 创建 ingress 默认 deny 规则
[root@qa-k8s-master01 networkpolicy]# cat ingress-def.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-all-ingress
spec:
  podSelector: {}
  policyTypes:
  - Ingress
```

```
// 应用到 dev 命名空间上
kubectl apply -f ingress-def.yaml -n dev
```

```
// 创建 ingress 默认放行规则
[root@qa-k8s-master01 networkpolicy]# cat ingress-def.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-all-ingress
spec:
  podSelector: {}
  ingress:
  - {}
  policyTypes:
  - Ingres
```

```
[root@qa-k8s-master01 networkpolicy]# cat allow-netpol-demo.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-myapp-ingress
spec:
  podSelector:
    matchLabels:
      app: myapp
  ingress:
  - from:
    - ipBlock:
        cidr: 10.244.0.0/16
        except:
        - 10.244.1.2/32
  ports:
  - protocol: TCP
    port: 80
```

calico的缺点：

路由的数目与容器数目相同，非常容易超过路由器，三层交换，甚至**node的处理能力**，从而限制了整个网络的扩张。

calico的每个**node**上会设置大量（海量）的**iptables**规则，路由，运维/排障难度大。

calico的原理决定了它不可能支持VPC，容器只能从calico设置的网段中获取ip。

calico目前的实现没有流量控制的功能，会出现少数容器抢占**node多数带宽的情况**。

calico的网络规模受到BGP网络规模的限制。

Dashboard

token

```
cd /etc/kubernetes/pki/
// 生成 dashboard 私钥
(umask 077; openssl genrsa -out dashboard.key 2048)
```

```
// 建立证书签署请求
openssl req -new -key dashboard.key -out dashboard.csr -subj "/O=mljr/CN=dashboard.mljr.com/"

// 签署证书
[root@qa-k8s-master01 pki]# openssl x509 -req -in dashboard.csr -CA ca.crt -CAkey ca.key -CAcreateserial -out dashboard.crt -days 3650
Signature ok
subject=/O=mljr/CN=dashboard.mljr.com
Getting CA Private Key

// 查看证书信息
[root@qa-k8s-master01 pki]# openssl x509 -in dashboard.crt -text -noout
Certificate:
Data:
    Version: 1 (0x0)
    Serial Number:
        c2:12:50:e1:21:65:5c:88
    Signature Algorithm: sha256WithRSAEncryption
    Issuer: CN=kubernetes
    Validity
        Not Before: Aug 19 04:01:52 2018 GMT
        Not After : Aug 16 04:01:52 2028 GMT
    Subject: O=mljr, CN=dashboard.mljr.com

// 创建 secret
kubectl create secret generic dashboard-cert -n kube-system --from-file=dashboard.crt=./dashboard.crt --from-file=dashboard.key=./dashboard.key

// 创建 serviceaccount
kubectl create clusterrolebinding dashboard-cluster-admin --clusterrole=cluster-admin --serviceaccount=kube-system:dashboard-admin

// 查看 token
kubectl get secret -n kube-system | grep dashboard-admin
kubectl describe secret dashboard-admin-token-vqdnr -n kube-system

// 部署 Dashboard
kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/master/src/deploy/recommended/kubernetes-dashboard.yaml

// 设置 service NodePort
kubectl patch svc kubernetes-dashboard -p '{"spec":{"type":"NodePort"}}' -n kube-system
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP, 53/TCP
5d				
kubernetes-dashboard	NodePort	10.97.156.106	<none>	443:31309/TCP
4m				

// token 登陆 dashboard
<https://172.28.28.70:31309/#!/login>

Kubernetes 仪表板

Kubeconfig
请选择您已配置用来访问集群的 kubeconfig 文件, 请浏览[配置对多个集群的访问](#)一节, 了解更多关于如何配置和使用 kubeconfig 文件的信息

令牌
每个服务帐号都有一条保密字典保存持有者令牌, 用来在仪表板登录, 请浏览[验证](#)一节, 了解更多关于如何配置和使用持有者令牌的信息

输入令牌
.....

登录 跳过

kubeconfig

```
# 实现 对 default 完全管理

// 创建 serviceaccount
kubectl create serviceaccount def-ns-default -n default
```

```
// 绑定
kubectl create rolebinding def-ns-admin --clusterrole=admin --serviceaccount=default:default

// 此时 kubectl get secret | grep def-ns-default 找出的账号 的 token 对 default
有完全管理权限
kubectl get secret def-ns-default-token-x4w2d -o jsonpath={.data.token}

kubectl config set-credentials dashboard-admin --
```

Monitor

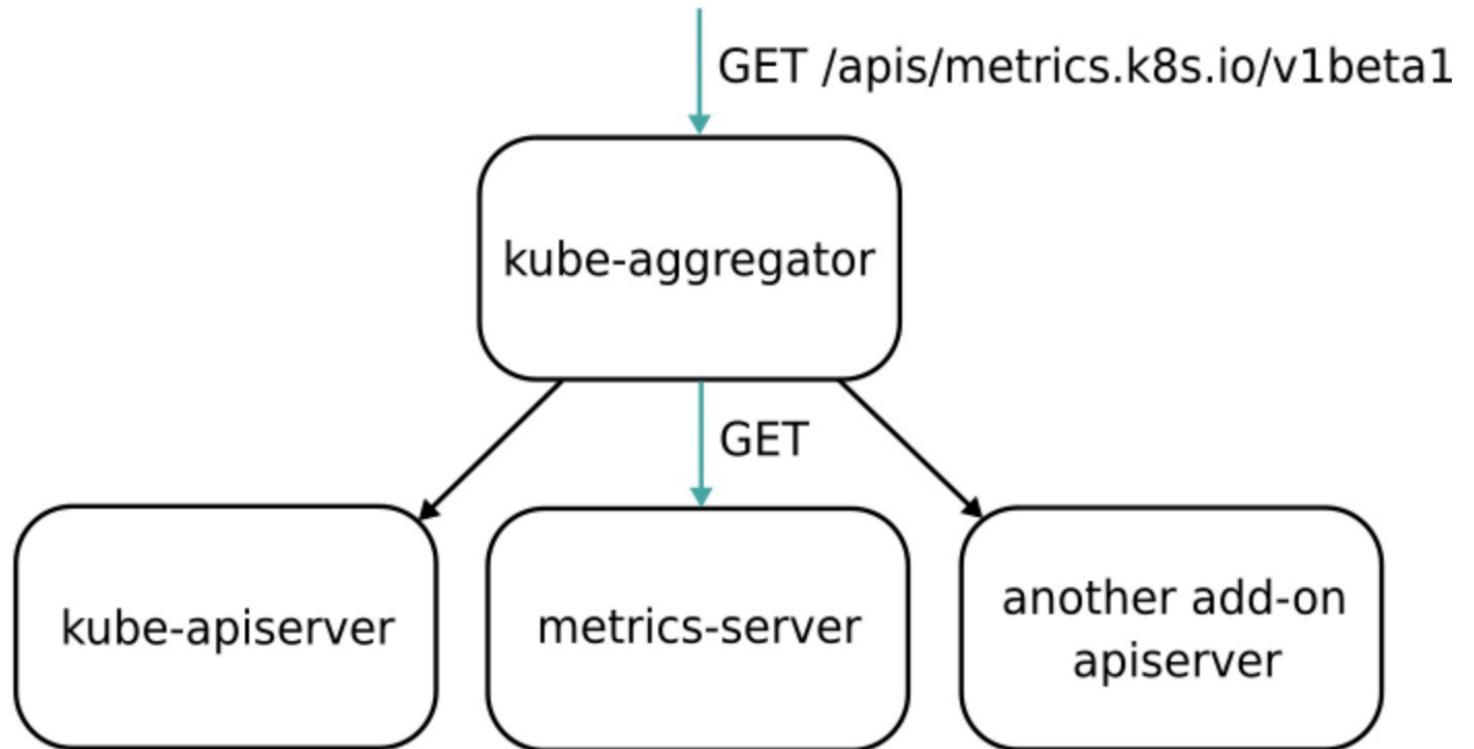
Metrics-Server

Kubernetes Metrics Server 是一个集群范围的资源使用数据聚合器，是 Heapster 的继承者。

The metrics server collects CPU and memory usage **for** nodes and pods **by** pooling data **from** the kubernetes.summary_api.

核心指标流水线：由kubelet, metrics-server以及由API **server**提供的api组成；
CPU累积使用率，内存实时使用率，
Pod的资源占用率及容器的磁盘占用率；

监控流水线：用于从系统收集各种指标数据并提供终端用户，存储系统以及HPA，它们包含核心指标及许多非核心指标。
非核心指标本身不能被k8s所解析。



// 部署 Metrics Server 服务后，`kubectl top` 命令就可以访问了

```
nohup kubectl proxy --port 8080 &
curl http://localhost:8080/apis/metrics.k8s.io/v1beta1/nodes
```

customresourcedefinitions

在 Kubernetes 中一切都可视为资源，Kubernetes 1.7 之后增加了对 CRD 自定义资源二次开发能力来扩展 Kubernetes API。

通过 CRD 我们可以向 Kubernetes API 中增加新资源类型，而不需要修改 Kubernetes 源码或创建自定义的 API server，该功能大大提高了 Kubernetes 的扩展能力。

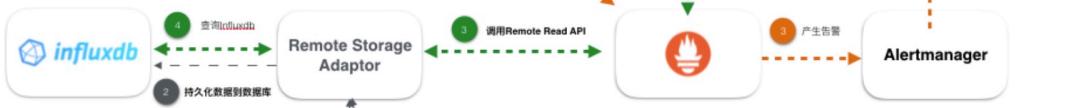
Prometheus

<https://github.com/kubernetes/kubernetes/tree/master/cluster/addons/prometheus>

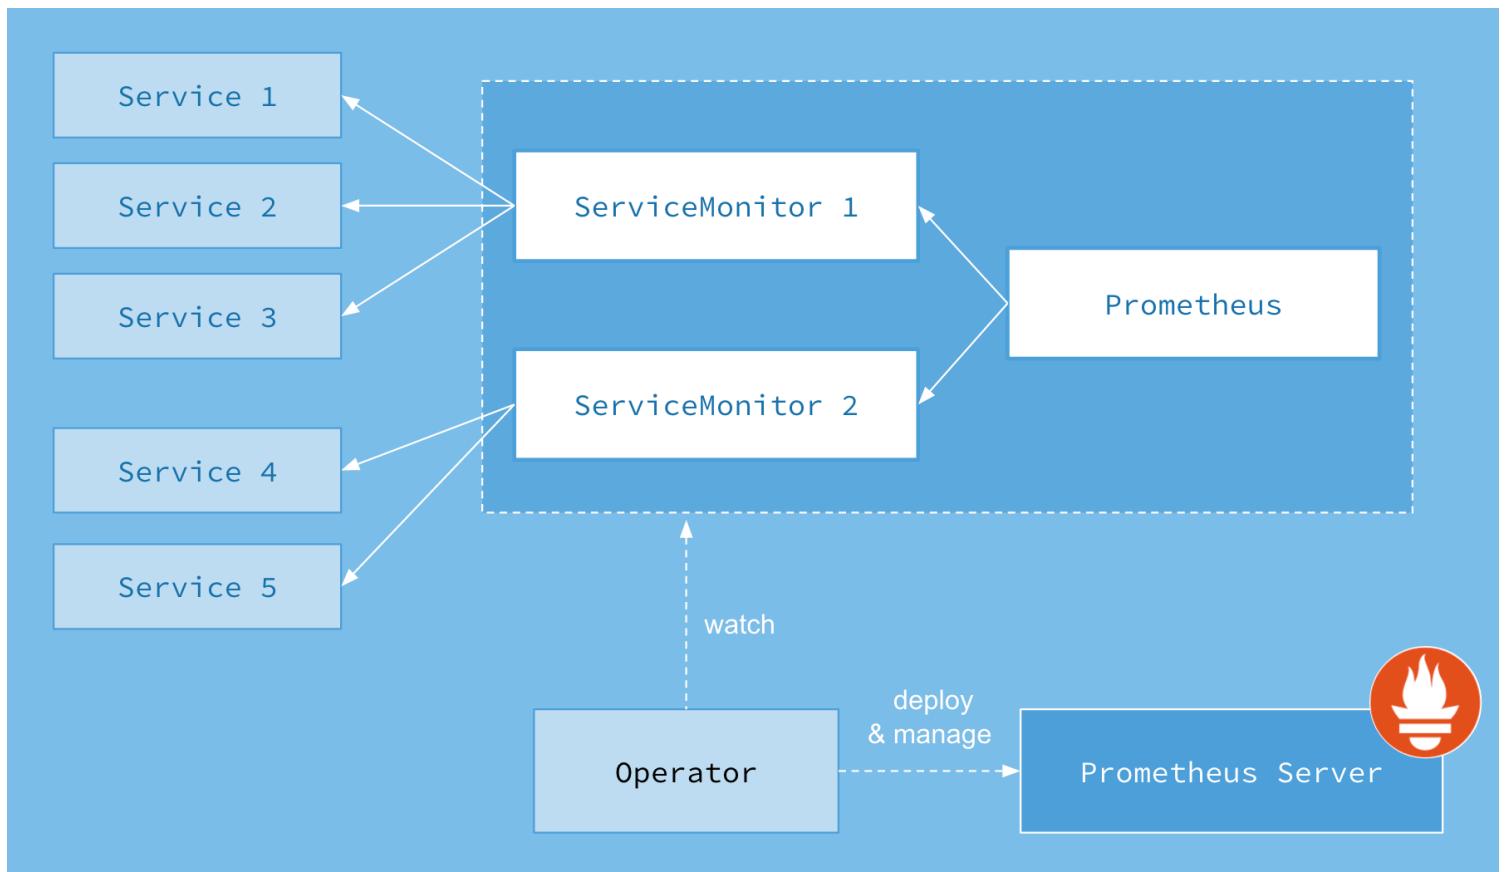
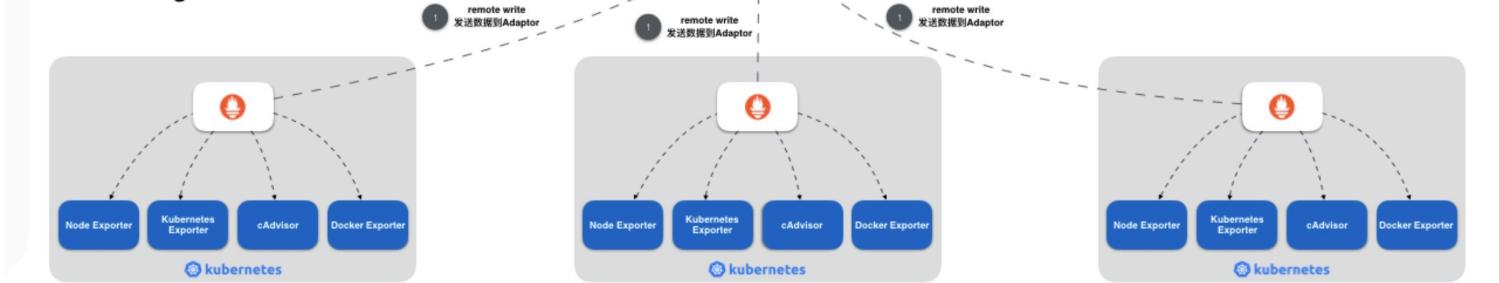
Client Panel



Prometheus Controller Panel



Prometheus Agent Panel



1) **Operator:** Operator 是整个系统的主要控制器。

Operator 以 Deployment 方式执行于 Kubernetes 集群上，並根据自定义资源(Custom Resource Definition, CRDs)来负责管理与部署 Prometheus Server。

同时监控这些自定义资源事件的变化来做相应的处理.

2) **Prometheus Server**: **Operator** 根据自定义资源 **Prometheus**类型中定义的内容而部署的 **Prometheus Server** 集群, 这些自定义资源可以看作是用来管理 **Prometheus Server** 集群的 **Stateful Sets** 资源.

3) **Prometheus**: **Prometheus**资源是声明性地描述 **Prometheus**部署的期望状态.

4) **ServiceMonitor**: **ServiceMonitor** 也是一个自定义资源, 它描述了一组被 **Prometheus** 监控的 **targets** 列表.

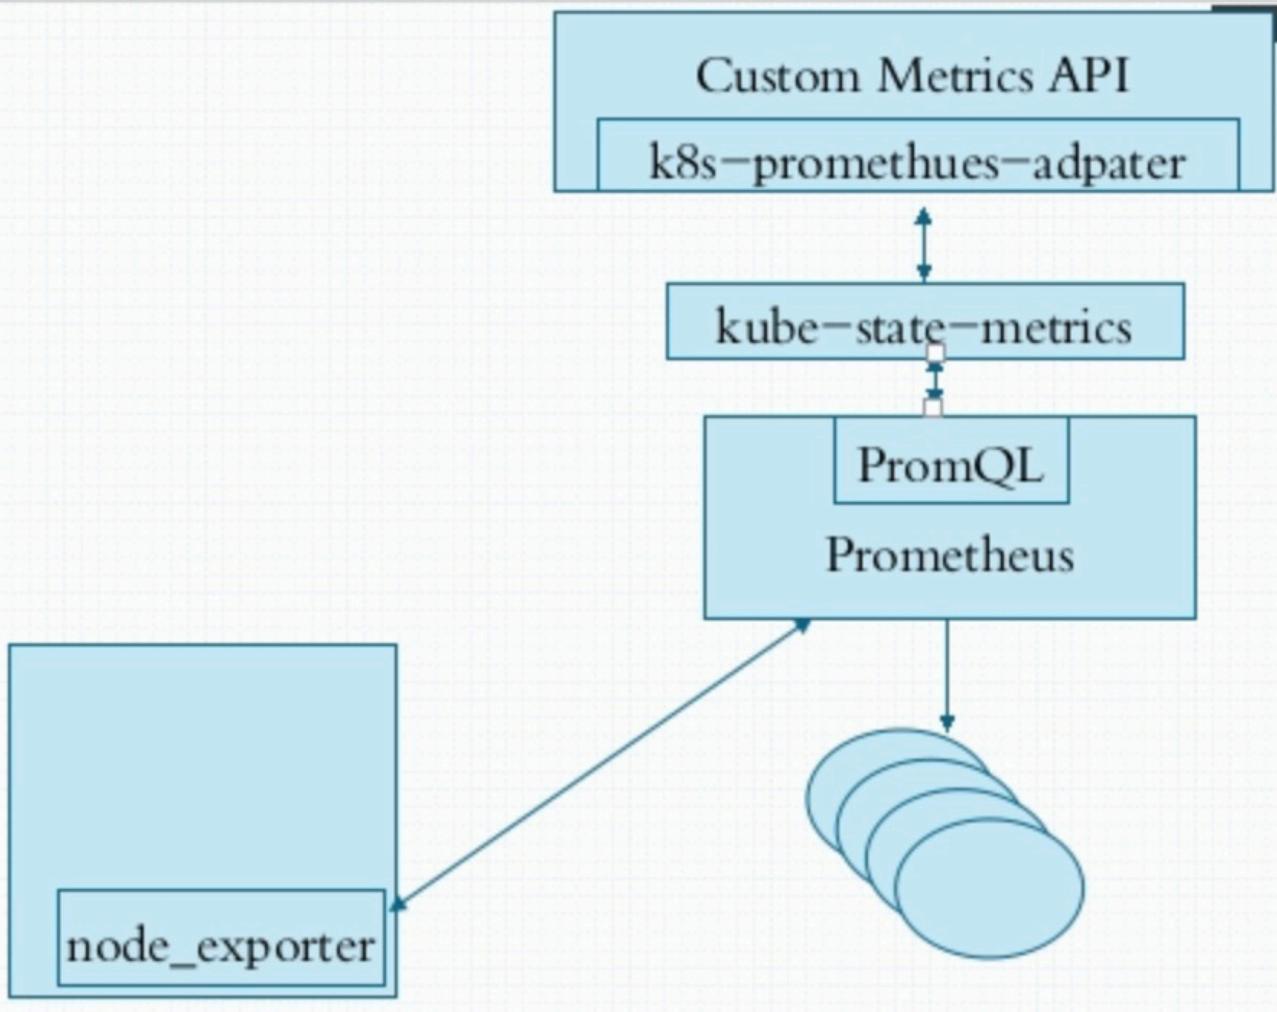
该资源通过 **Labels** 来选取对应的 **Service Endpoint**, 让 **Prometheus Server** 通过选取的 **Service** 来获取 **Metrics** 信息.

5) **Service**: **Service** 资源主要用来对应 **Kubernetes** 集群中的 **Metrics Server Pod**, 来提供给 **ServiceMonitor** 选取让 **Prometheus Server** 来获取信息.

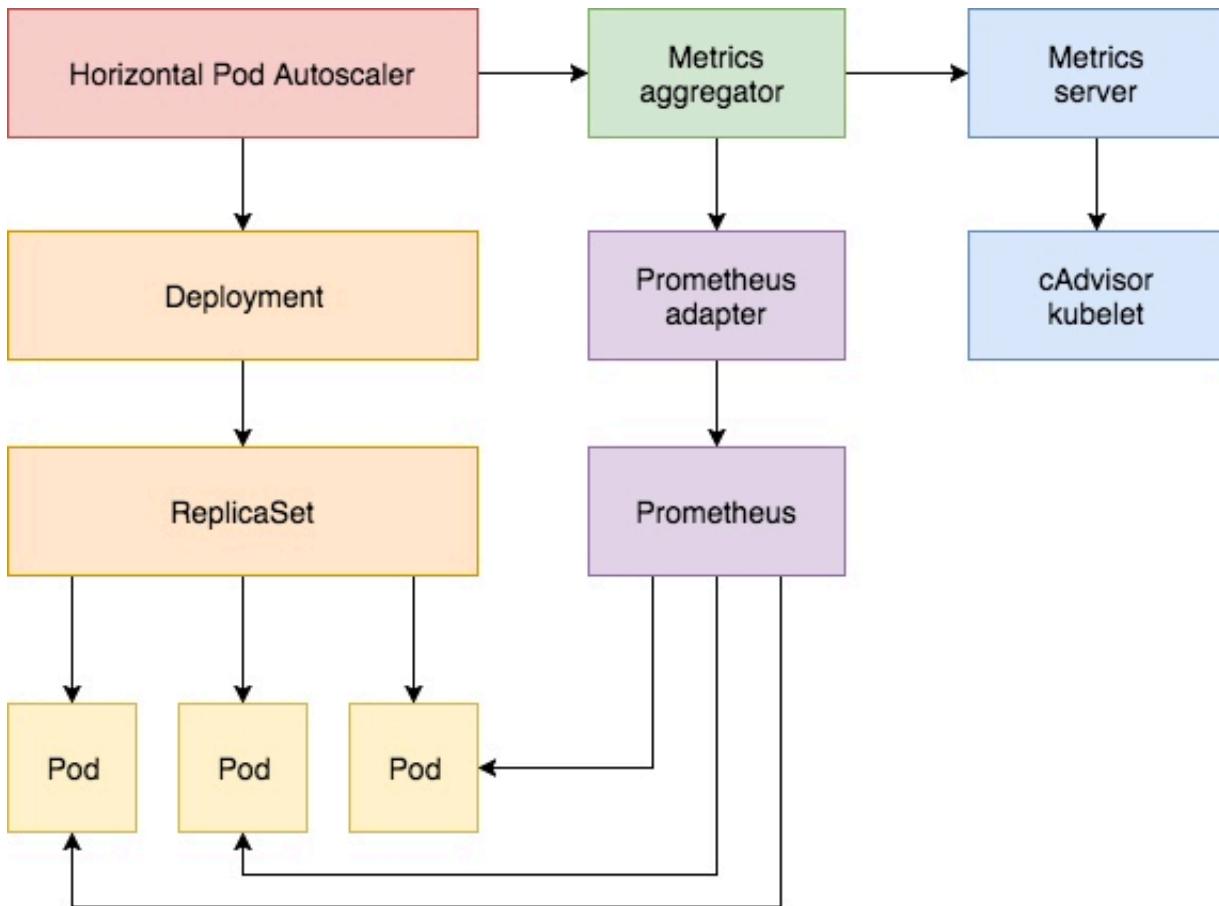
简单的说就是 **Prometheus** 监控的对象, 例如之前了解的 **Node Exporter Service**, **Mysql Exporter Service** 等等.

6) **Alertmanager**: **Alertmanager** 也是一个自定义资源类型, 由 **Operator** 根据资源描述内容来部署 **Alertmanager** 集群.

<https://github.com/camilb/prometheus-kubernetes>

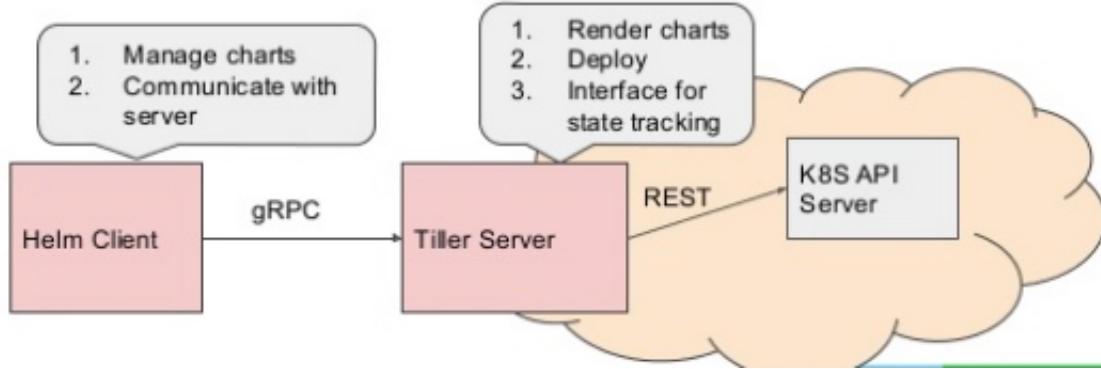


HPA



helm

Basic Architecture



Charts: A bundle of kubernetes resource data, it operates the System Package

Repository: A collection of release charts, something like NPM, or Ubuntu Repository.

Release: A chart instance that was loaded into Kubernetes. If the same chart is installed several times inside the same cluster, each time the cart will have its own release.

```
 wget https://storage.googleapis.com/kubernetes-helm/helm-v2.9.1-linux-amd64.tar.gz  
 tar xf helm-v2.9.1-linux-amd64.tar.gz  
 cp linux-amd64/helm /usr/local/bin/
```

核心术语：

Chart:一个helm程序包；

Repository:Charts仓库, https/http服务器；

Release:特定的Chart部署于目标集群上的一个实例；

Chart -> Config -> Release

程序架构：

helm:客户端, 管理本地的Chart仓库, 管理Chart, 与Tiller服务器交互, 发送Chart, 实例安装、查询、卸载等

Tiller:服务端, 接收helm发来的Charts与Config, 合并生成release;

RBAC配置文件示例: <https://github.com/helm/helm/blob/master/docs/rbac.md>

官方可用的Chart列表: <https://hub.kubeapps.com/>

tiller

```
// tiller rbac  
  
kubectl apply -f tiller-rbac-config.yaml  
  
// tiller init  
helm init --service-account tiller
```

```
[root@qa-k8s-master01 helm]# helm version  
Client: &version.Version{SemVer:"v2.9.1", GitCommit:"20adb27c7c5868466912eebdf6  
664e7390ebe710", GitTreeState:"clean"}  
Server: &version.Version{SemVer:"v2.9.1", GitCommit:"20adb27c7c5868466912eebdf6  
664e7390ebe710", GitTreeState:"clean"}  
[root@qa-k8s-master01 helm]# helm repo list  
NAME      URL  
stable    https://kubernetes-charts.storage.googleapis.com  
local     http://127.0.0.1:8879/charts
```

Gitlab Install By Helm

```
// 安装
helm install --name gitlab -f values.yaml stable/gitlab-ce

// 更新
helm upgrade -f values.yaml gitlab stable/gitlab-ce
```

Creating Highly Available Clusters with kubeadm

```
172.28.28.75 qa-k8s-master01.mljr.com master01
172.28.28.71 qa-k8s-master02.mljr.com master02
172.28.28.76 qa-k8s-master03.mljr.com master03
172.28.28.72 qa-k8s-node01.mljr.com node01
172.28.28.73 qa-k8s-node02.mljr.com node02
172.28.28.74 qa-k8s-node03.mljr.com node03
172.28.3.55 qa-k8s-node04.mljr.com node04
```

Init

```
// 在所有的 master 节点上 执行
yum install -y kubelet kubeadm kubectl docker-ce
systemctl enable docker && systemctl start docker
systemctl enable kubelet

// 更多 init 需要看 单 master 集群内容

curl -o /usr/local/bin/cfssl https://pkg.cfssl.org/R1.2/cfssl_linux-amd64
curl -o /usr/local/bin/cfssljson https://pkg.cfssl.org/R1.2/cfssljson_linux-amd
64
chmod +x /usr/local/bin/cfssl*
```

Create load balancer for kube-apiserver

master01

```
yum install haproxy -y

cd /etc/haproxy/
mv haproxy.cfg haproxy.cfg.orgin

cat >/etc/haproxy/haproxy.cfg<<EOF
global
    log 127.0.0.1 local0
    log 127.0.0.1 local1 notice
    tune.ssl.default-dh-param 2048

defaults
    log global
    mode http
    option dontlognull
    timeout connect 5000ms
    timeout client 600000ms
    timeout server 600000ms

listen stats
    bind :9090
    mode http
    balance
    stats uri /haproxy_stats
    stats auth admin:admin123
    stats admin if TRUE

frontend kube-apiserver-https
    mode tcp
    bind :8443
    default_backend kube-apiserver-backend

backend kube-apiserver-backend
    mode tcp
    balance roundrobin
    stick-table type ip size 200k expire 30m
    stick on src
    server apiserver1 172.28.28.75:6443 check
    server apiserver2 172.28.28.71:6443 check
    server apiserver3 172.28.28.76:6443 check
EOF

systemctl start haproxy
systemctl enable haproxy
```

```
yum install keepalived -y
/etc/keepalived

mv keepalived.conf keepalived.conf.origin

cat > keepalived.conf <<EOF
! Configuration File for keepalived
global_defs {
    router_id LVS_DEVEL
}

vrrp_instance VI_1 {
    state MASTER
    interface ens160
    virtual_router_id 51
    priority 90
    authentication {
        auth_type PASS
        auth_pass 35b4c9t019x4d1600c483e10ad1d
    }
    virtual_ipaddress {
        172.28.28.200
    }
}
EOF

systemctl start keepalived
systemctl enable keepalived
```

master02 / master03

```
// haproxy 配置文件同 master01

// keepalived 配置文件
state => BCKUP
priority => < 90
```

Etcd cluster

先在 master01 生成 etcd 证书

```
mkdir -p /etc/kubernetes/pki/etcd
cd /etc/kubernetes/pki/etcd

cat >ca-config.json <<EOL
{
    "signing": {
        "default": {
            "expiry": "87600h"
        },
        "profiles": {
            "server": {
                "expiry": "87600h",
                "usages": [
                    "signing",
                    "key encipherment",
                    "server auth",
                    "client auth"
                ]
            },
            "client": {
                "expiry": "87600h",
                "usages": [
                    "signing",
                    "key encipherment",
                    "client auth"
                ]
            },
            "peer": {
                "expiry": "87600h",
                "usages": [
                    "signing",
                    "key encipherment",
                    "server auth",
                    "client auth"
                ]
            }
        }
    }
}
EOL
```

```
cat >ca-csr.json <<EOL
{
    "CN": "etcd",
```

```

    "key": {
        "algo": "rsa",
        "size": 2048
    }
}
EOL

// 生成 CA 证书
cfssl gencert -initca ca-csr.json | cfssljson -bare ca -

cat >client.json <<EOL
{
    "CN": "client",
    "key": {
        "algo": "ecdsa",
        "size": 256
    }
}
EOL

cfssl gencert -ca=ca.pem -ca-key=ca-key.pem -config=ca-config.json -profile=client client.json | cfssljson -bare client

scp -r -P 37100 /etc/kubernetes/pki/etcd master02:/etc/kubernetes/pki/
scp -r -P 37100 /etc/kubernetes/pki/etcd master03:/etc/kubernetes/pki/

```

在所有的 master 都执行

```

mkdir -p /etc/kubernetes/pki/etcd
cd /etc/kubernetes/pki/etcd

export PEER_NAME=$(hostname)
export PRIVATE_IP=$(ip addr show ens160 | grep -Po 'inet \K[\d.]+')

cfssl print-defaults csr > config.json
sed -i '0,/CN/{s/example\.net/"$PEER_NAME"/}' config.json
sed -i 's/www\.example\.net/"$PRIVATE_IP"/' config.json
sed -i 's/example\.net/"$PUBLIC_IP"/' config.json

cfssl gencert -ca=ca.pem -ca-key=ca-key.pem -config=ca-config.json -profile=server config.json | cfssljson -bare server
cfssl gencert -ca=ca.pem -ca-key=ca-key.pem -config=ca-config.json -profile=pee

```

```
r config.json | cfssljson -bare peer

cd /usr/local/src
export ETCD_VERSION=v3.3.9
curl -SSL https://github.com/coreos/etcd/releases/download/${ETCD_VERSION}/etcd-$ETCD_VERSION-linux-amd64.tar.gz | tar -xv --strip-components=1 -C /usr/local/bin/
rm -rf etcd-$ETCD_VERSION-linux-amd64*

touch /etc/etcd.env
echo "PEER_NAME=$PEER_NAME" >> /etc/etcd.env
echo "PRIVATE_IP=$PRIVATE_IP" >> /etc/etcd.env

cat >/etc/systemd/system/etcd.service <<EOL
[Unit]
Description=etcd
Documentation=https://github.com/coreos/etcd
Conflicts=etcd.service
Conflicts=etcd2.service

[Service]
EnvironmentFile=/etc/etcd.env
Type=notify
Restart=always
RestartSec=5s
LimitNOFILE=40000
TimeoutStartSec=0

ExecStart=/usr/local/bin/etcd --name ${PEER_NAME} \
--data-dir /var/lib/etcd \
--listen-client-urls https://${PRIVATE_IP}:2379 \
--advertise-client-urls https://${PRIVATE_IP}:2379 \
--listen-peer-urls https://${PRIVATE_IP}:2380 \
--initial-advertise-peer-urls https://${PRIVATE_IP}:2380 \
--cert-file=/etc/kubernetes/pki/etcd/server.pem \
--key-file=/etc/kubernetes/pki/etcd/server-key.pem \
--client-cert-auth \
--trusted-ca-file=/etc/kubernetes/pki/etcd/ca.pem \
--peer-cert-file=/etc/kubernetes/pki/etcd/peer.pem \
--peer-key-file=/etc/kubernetes/pki/etcd/peer-key.pem \
--peer-client-cert-auth \
--peer-trusted-ca-file=/etc/kubernetes/pki/etcd/ca.pem \
--initial-cluster qa-k8s-master01.mljr.com=https://172.28.28.75:2380,qa-k8s-master02.mljr.com=https://172.28.28.71:2380,qa-k8s-master03.mljr.com=https://172.28.28.76:2380 \

```

```
--initial-cluster-token mljr-etcd-token |  
--initial-cluster-state new
```

[Install]

```
WantedBy=multi-user.target  
EOL
```

```
systemctl daemon-reload  
systemctl start etcd  
systemctl enable etcd
```

检查集群状态

```
[root@qa-k8s-master01 etcd]# etcdctl --cert-file /etc/kubernetes/pki/etcd/peer.pem --key-file /etc/kubernetes/pki/etcd/peer-key.pem --ca-file /etc/kubernetes/pki/etcd/ca.pem --endpoints https://172.28.28.71:2379,https://172.28.28.76:2379  
cluster-health  
member 555d61d38df58f48 is healthy: got healthy result from https://172.28.28.75:2379  
member c494646ef59eff60 is healthy: got healthy result from https://172.28.28.71:2379  
member eb5d91cdcad637f0 is healthy: got healthy result from https://172.28.28.76:2379  
cluster is healthy
```

```
[root@qa-k8s-master01 etcd]# etcdctl --cert-file /etc/kubernetes/pki/etcd/peer.pem --key-file /etc/kubernetes/pki/etcd/peer-key.pem --ca-file /etc/kubernetes/pki/etcd/ca.pem --endpoints https://172.28.28.71:2379,https://172.28.28.76:2379  
member list  
555d61d38df58f48: name=qa-k8s-master01.mljr.com peerURLs=https://172.28.28.75:2380 clientURLs=https://172.28.28.75:2379 isLeader=false  
c494646ef59eff60: name=qa-k8s-master02.mljr.com peerURLs=https://172.28.28.71:2380 clientURLs=https://172.28.28.71:2379 isLeader=false  
eb5d91cdcad637f0: name=qa-k8s-master03.mljr.com peerURLs=https://172.28.28.76:2380 clientURLs=https://172.28.28.76:2379 isLeader=true
```

创建 k8s 集群

在所有的 master 节点生成kubeadm-config.yaml文件

```
cat > kubeadm-config.yaml <<EOL
apiVersion: kubeadm.k8s.io/v1alpha2
kind: MasterConfiguration
kubernetesVersion: v1.11.2
apiServerCertSANs:
- "172.28.28.200"
api:
  controlPlaneEndpoint: "172.28.28.200:8443"
etcd:
  external:
    endpoints:
    - https://172.28.28.75:2379
    - https://172.28.28.71:2379
    - https://172.28.28.76:2379
    caFile: /etc/kubernetes/pki/etcd/ca.pem
    certFile: /etc/kubernetes/pki/etcd/client.pem
    keyFile: /etc/kubernetes/pki/etcd/client-key.pem
networking:
  # This CIDR is a calico default. Substitute or remove for your CNI provider
  -
    podSubnet: "10.244.0.0/16"
EOL
```

在 master01 上执行

```
kubeadm init --config kubeadm-config.yaml

// 执行成功后 保留添加 node 命令
kubeadm join 172.28.28.200:8443 --token 99fxwc.wbjgtzf10z2sn1d2 --discovery-token-ca-cert-hash sha256:31fcbe3cc4da9bb872c2d0fab6ac4747c0634246796060dfa0d8265bf9bfbc76

// 配置本地命令行管理
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

// cp ca信息到 master02 master03
scp -P 37100 /etc/kubernetes/pki/* master02:/etc/kubernetes/pki/
scp -P 37100 /etc/kubernetes/pki/* master03:/etc/kubernetes/pki/
```

在 master02 和 master03 上执行 init

```
kubeadm init --config kubeadm-config.yaml
```

Install CNI

Flannel And Enable Directrouting

见 CNI

Installing Calico for policy and flannel for networking

见 CNI

Add Node

// 同单节点，略

Dashboard

部署Web UI

```
wget https://raw.githubusercontent.com/kubernetes/dashboard/master/src/deploy/recommended/kubernetes-dashboard.yaml
```

下载并编辑kubernetes-dashboard.yaml，并在此文件中的Service部分下添加`type: NodePort`和`nodePort: 30443`

修改部分如下。

```
kind: Service
apiVersion: v1
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
```

```

namespace: kube-system
spec:
  ports:
    - port: 443
      targetPort: 8443
      nodePort: 30443
    type: "NodePort"
  selector:
    k8s-app: kubernetes-dashboard

```

```
kubectl apply -f kubernetes-dashboard.yaml
```

// 访问

使用集群的任何一个 ip 都能访问

<https://172.28.28.71:30443>

```
[root@qa-k8s-master01 dashboard]# kubectl get svc -n kube-system
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)
AGE
kube-dns       ClusterIP 10.96.0.10  <none>        53/UDP,53/TCP
  4h
kubernetes-dashboard   NodePort  10.103.31.215 <none>        443:30443/TCP
  3m
```

创建一个管理员用户

```

cat >admin-user.yaml <<EOF
#----Create Service Account----#
apiVersion: v1
kind: ServiceAccount
metadata:
  name: admin-user
  namespace: kube-system
---#
#----Create ClusterRoleBinding----#
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  name: admin-user

```

```
roleRef:  
  apiGroup: rbac.authorization.k8s.io  
  kind: ClusterRole  
  name: cluster-admin  
subjects:  
- kind: ServiceAccount  
  name: admin-user  
  namespace: kube-system  
EOF
```

token

```
[root@qa-k8s-master01 dashboard]# kubectl apply -f amind-user.yaml  
serviceaccount/admin-user created  
clusterrolebinding.rbac.authorization.k8s.io/admin-user created
```

// 查看 token 相关信息

```
[root@qa-k8s-master01 dashboard]# kubectl -n kube-system describe secret $(kubectl -n kube-system get secret | grep admin-user | awk '{print $1}')  
Name:           admin-user-token-fxls7  
Namespace:      kube-system  
Labels:         <none>  
Annotations:    kubernetes.io/service-account.name=admin-user  
                kubernetes.io/service-account.uid=93350806-ab7c-11e8-9192-0050569d46cc  
d46cc
```

Type: kubernetes.io/service-account-token

Data

====

```
ca.crt:      1025 bytes  
namespace:   11 bytes  
token:       eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcm5ldGVzL3Nlc...  
VhY2NvdW50Iiwiia3ViZXJuZXRLcy5pb...  
3RlbSI...  
LXRva2VuLWZ4bHM3Ii...  
ubmFtZSI6ImFkbWluLXVzZ...  
Njb3VudC51aWQ...  
XN0ZW06c2VydmljZWFjY291bnQ6a3ViZS1zeXN0ZW06YW...  
wMqQDkjFKDR8R3TSIIcy_SKoEau-j1yFL2huE89w2y5UDVV8DBoUmCXrHEc8acps0krFAfSqeSK1kkX  
WCxz...  
4iWGP07QyzpLh2vn7BAYv...  
0lfkQP0ylBr0WnRa19T0zP2VgfLeHdCyH4TWLgh-0RGxrF8rBufD00cVNQ5teP0C_hl00k1g-urNuw7  
zR0r4lQ
```

```
// 命令行获取 token
[root@qa-k8s-master01 dashboard]# kubectl -n kube-system get secret $(kubectl -n kube-system get secret | grep admin-user | awk '{print $1}') -o jsonpath={.data.token} | awk '{print $0}' | base64 -d | awk '{print $0}'
```

kubeconfig Login

添加一个用户对 dev namespace 有管理权限

```
[root@qa-k8s-master01 dashboard]# cat > dev-admin.yaml <<EOF
#----Create dev NS----#
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: dev
spec: {}
status: {}

---
#----Create Rule----#
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  namespace: dev
  name: dev-admin
rules:
- apiGroups: ["*"]
  resources: ["*"]
  verbs: ["*"]

---
#----Create Service Account----#
apiVersion: v1
kind: ServiceAccount
metadata:
  name: dev-admin
  namespace: dev
```

```

---
#-----Create RoleBinding-----#
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: dev-admin
  namespace: dev
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: dev-admin
subjects:
- kind: ServiceAccount
  name: dev-admin
  namespace: dev
EOF
[root@qa-k8s-master01 dashboard]# kubectl apply -f default-admin.yaml

```

```

dev_token=`kubectl -n dev get secret $(kubectl -n dev get secret | grep dev-admin | awk '{print $1}') --output yaml | jq '.data.token' | base64 -d | awk '{print $0}'` 

kubectl config set-credentials dev-admin --token=$dev_token --kubeconfig=/root/dev-ns-admin.conf

kubectl config set-context dev-admin@kubernetes --cluster=kubernetes --user=dev-admin --kubeconfig=/root/dev-ns-admin.conf

kubectl config use-context dev-admin@kubernetes --kubeconfig=/root/dev-ns-admin.conf

[root@qa-k8s-master01 dashboard]# kubectl config view --kubeconfig=/root/dev-ns-admin.conf

```

Error

```

r: code = Unknown desc = failed to set up sandbox container "ab14be8f0b48977fae b57576a757997e9dc1efaedf394d2ec1162713118d49a0" network for pod "coredns-78fcdf6894-6872b": NetworkPlugin cnetwork failed to set up pod "coredns-78fcdf6894-6872b_kube-system" network: failed to set bridge addr: "cni0" already has an IP address

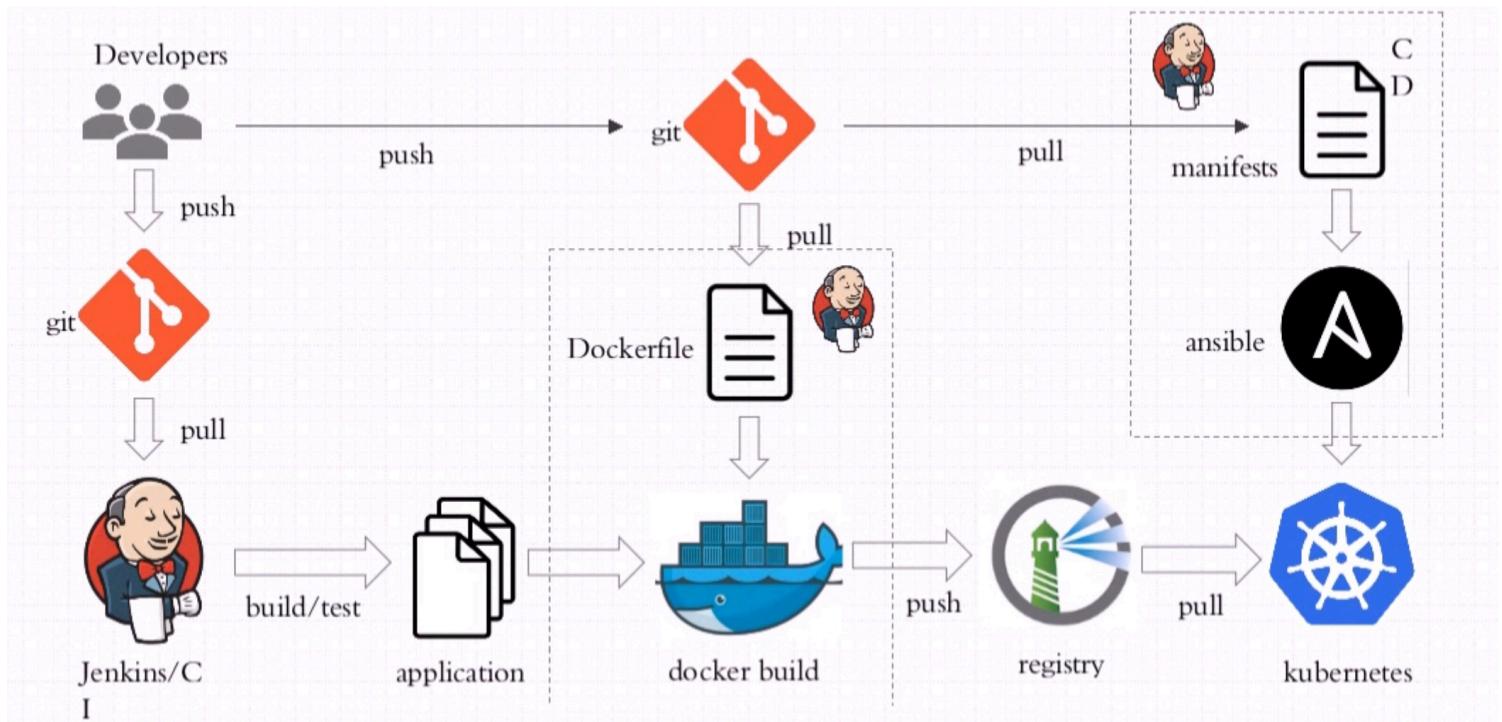
```

s different from 10.244.0.1/24

==> 此时coredns 容器是运行不起来的 出现这种情况是因为之前有些机器已经做 node 了，纯净服务器不会出现这样问题。

```
systemctl stop kubelet  
ifconfig cni0 down  
brctl delbr cni0  
ip link delete flannel.1  
systemctl start kubelet  
reboot
```

K8s With Paas



Ingress

Ingress 使用开源的反向代理负载均衡器来实现对外暴漏服务。

启用 Ingress

gitlab 实例

```
// 准备
```

在之前集群的基础上，新添加一个节点 **node5**(生成环境建议3个以上)，并设置污点，专用于暴露服务给外部.

使用 **Traefik(80)** 和 **Ingress_nginx(8080)** 两种方式提供服务.

基于https, **Traefik(443)** 和 **Ingress_nginx(8443)**

hostNetwork

新加 node 并设置污点

gitlab安装

```
// 需要创建4个 pv，具体大小看要求.
```

```
# cat pv-demo.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv002
  labels:
    name: pv002
spec:
  nfs:
    path: /pv/volumes/v2
    server: master03
  accessModes: ["ReadWriteMany", "ReadWriteOnce"]
  capacity:
    storage: 20Gi
```

```
[root@qa-k8s-master01 gitlab]# cat values.yaml
## GitLab CE image
## ref: https://hub.docker.com/r/gitlab/gitlab-ce/tags/
##
image: gitlab/gitlab-ce:9.4.1-ce.0
```

```
## Specify a imagePullPolicy
## 'Always' if imageTag is 'latest', else set to 'IfNotPresent'
## ref: http://kubernetes.io/docs/user-guide/images/#pre-pulling-images
##
# imagePullPolicy:
imagePullPolicy: IfNotPresent

## The URL (with protocol) that your users will use to reach the install.
## ref: https://docs.gitlab.com/omnibus/settings/configuration.html#configuring
## the external url for gitlab
##
#externalUrl: http://your-domain.com/
externalUrl: http://my.gitlab.com/

## Change the initial default admin password if set. If not set, you'll be
## able to set it when you first visit your install.
##
#gitlabRootPassword: ""
gitlabRootPassword: "ss1234560.0"

## For minikube, set this to NodePort, elsewhere use LoadBalancer
## ref: http://kubernetes.io/docs/user-guide/services/#publishing-services---se
## rvice-types
##
#serviceType: LoadBalancer
serviceType: ClusterIP

## Ingress configuration options
##
ingress:
  annotations:
    # kubernetes.io/ingress.class: nginx
    # kubernetes.io/tls-acme: "true"
  enabled: false
  tls:
    # - secretName: gitlab.cluster.local
    #   hosts:
    #     - gitlab.cluster.local
  url: gitlab.cluster.local

## Configure external service ports
## ref: http://kubernetes.io/docs/user-guide/services/
sshPort: 22
httpPort: 80
```

```
httpsPort: 443
## livenessPort Port of liveness probe endpoint
livenessPort: http
## readinessPort Port of readiness probe endpoint
readinessPort: http

## Configure resource requests and limits
## ref: http://kubernetes.io/docs/user-guide/compute-resources/
##
resources:
  ## GitLab requires a good deal of resources. We have split out Postgres and
  ## redis, which helps some. Refer to the guidelines for larger installs.
  ## ref: https://docs.gitlab.com/ce/install/requirements.html#hardware-requirements
  requests:
    memory: 1Gi
    cpu: 500m
  limits:
    memory: 2Gi
    cpu: 1

## Enable persistence using Persistent Volume Claims
## ref: http://kubernetes.io/docs/user-guide/persistent-volumes/
## ref: https://docs.gitlab.com/ce/install/requirements.html#storage
##
persistence:
  ## This volume persists generated configuration files, keys, and certs.
  ##
  gitlabEtc:
    enabled: true
    size: 1Gi
    ## If defined, volume.beta.kubernetes.io/storage-class: <storageClass>
    ## Default: volume.alpha.kubernetes.io/storage-class: default
    ##
    # storageClass:
    accessMode: ReadWriteOnce
    ## This volume is used to store git data and other project files.
    ## ref: https://docs.gitlab.com/omnibus/settings/configuration.html#storing-git-data-in-an-alternative-directory
    ##
  gitlabData:
    enabled: true
    size: 10Gi
    ## If defined, volume.beta.kubernetes.io/storage-class: <storageClass>
    ## Default: volume.alpha.kubernetes.io/storage-class: default
```

```
##  
# storageClass:  
accessMode: ReadWriteOnce  
  
## Configuration values for the postgresql dependency.  
## ref: https://github.com/kubernetes/charts/blob/master/stable/postgresql/README.md  
##  
#  
postgresql:  
  # 9.6 is the newest supported version for the GitLab container  
  imageTag: "9.6"  
  cpu: 1000m  
  memory: 1Gi  
  
  postgresUser: gitlab  
  postgresPassword: gitlab  
  postgresDatabase: gitlab  
  
  persistence:  
    size: 10Gi  
  
## Configuration values for the redis dependency.  
## ref: https://github.com/kubernetes/charts/blob/master/stable/redis/README.md  
##  
redis:  
  redisPassword: "gitlab"  
  
  resources:  
    requests:  
      memory: 1Gi  
  
  persistence:  
    size: 10Gi
```

```
helm install --name gitlab -f values.yaml stable/gitlab-ce
```

Traefik 设置

```
// Traefik 略
```

准备 Traefik Ingress 配置文件

```
[root@qa-k8s-master01 Ingress]# cat ingress-traefik.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-traefik
  annotations:
    kubernetes.io/ingress.class: traefik
    traefik.frontend.rule.type: PathPrefixStrip
spec:
  rules:
  - host: my.gitlab.com
    http:
      paths:
      - path: /
        backend:
          serviceName: gitlab-gitlab-ce
          servicePort: 80
```

设置traefik service

```
[root@qa-k8s-master01 gitlab]# cat my-gitlab-com-traefik-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: ingress-traefik-gitlabe
  namespace: default

spec:
  type: NodePort
  ports:
  - name: http
    port: 80
    targetPort: 80
    nodePort: 30081
    protocol: TCP
  - name: https
    port: 443
    targetPort: 443
    protocol: TCP
  selector:
    app: gitlab-gitlab-ce
```

```
// 配置  
http://my.gitlab.com:30081/
```

Ingress_nginx 设置

```
// 安装 Ingress Controller 见上
```

准备 Nginx Ingress

```
[root@qa-k8s-master01 gitlab]# cat my-gitlab-com-ingress-nginx.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-gitlab-com
  namespace: default
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  rules:
  - host: my.gitlab.com
    http:
      paths:
      - backend:
          serviceName: gitlab-gitlab-ce
          servicePort: 80
```

准备 Nginx Ingress Service

```
// cat my-gitlab-com-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: ingress-nginx
  namespace: ingress-nginx
  labels:
    app.kubernetes.io/name: ingress-nginx
    app.kubernetes.io/part-of: ingress-nginx
spec:
```

```
type: NodePort
ports:
- name: http
  port: 80
  targetPort: 80
  nodePort: 30080
  protocol: TCP
- name: https
  port: 443
  targetPort: 443
  protocol: TCP
selector:
  app.kubernetes.io/name: ingress-nginx
```

访问 <http://my.gitlab.com:30080/>

https

```
(umask 077; openssl genrsa -out gitlabe.key 2048)
openssl req -new -x509 -key gitlabe.key -out gitlabe.crt -subj /C=CN/ST=Shenzh
eng/O=DevOps/CN=my.gitlab.com
kubectl create secret tls gitlabe-ingress-secret --cert=gitlabe.crt --key=gitla
be.key
```